

EMBEDDED MACHINE LEARNING

LECTURE 11 - ADVANCED NEURAL ARCHITECTURES

Holger Fröning

holger.froening@ziti.uni-heidelberg.de

HAWAI Group (formerly Computing Systems), Institute of Computer Engineering
Heidelberg University

WHY ALWAYS IMAGE CLASSIFICATION?

Image classification is a task that is very handy when it comes to feature extraction

Feature extraction is foundational for many other tasks

- Object detection & object tracking

- Image segmentation - semantic (cannot distinguish between different instances of the same category) or instance (can distinguish)

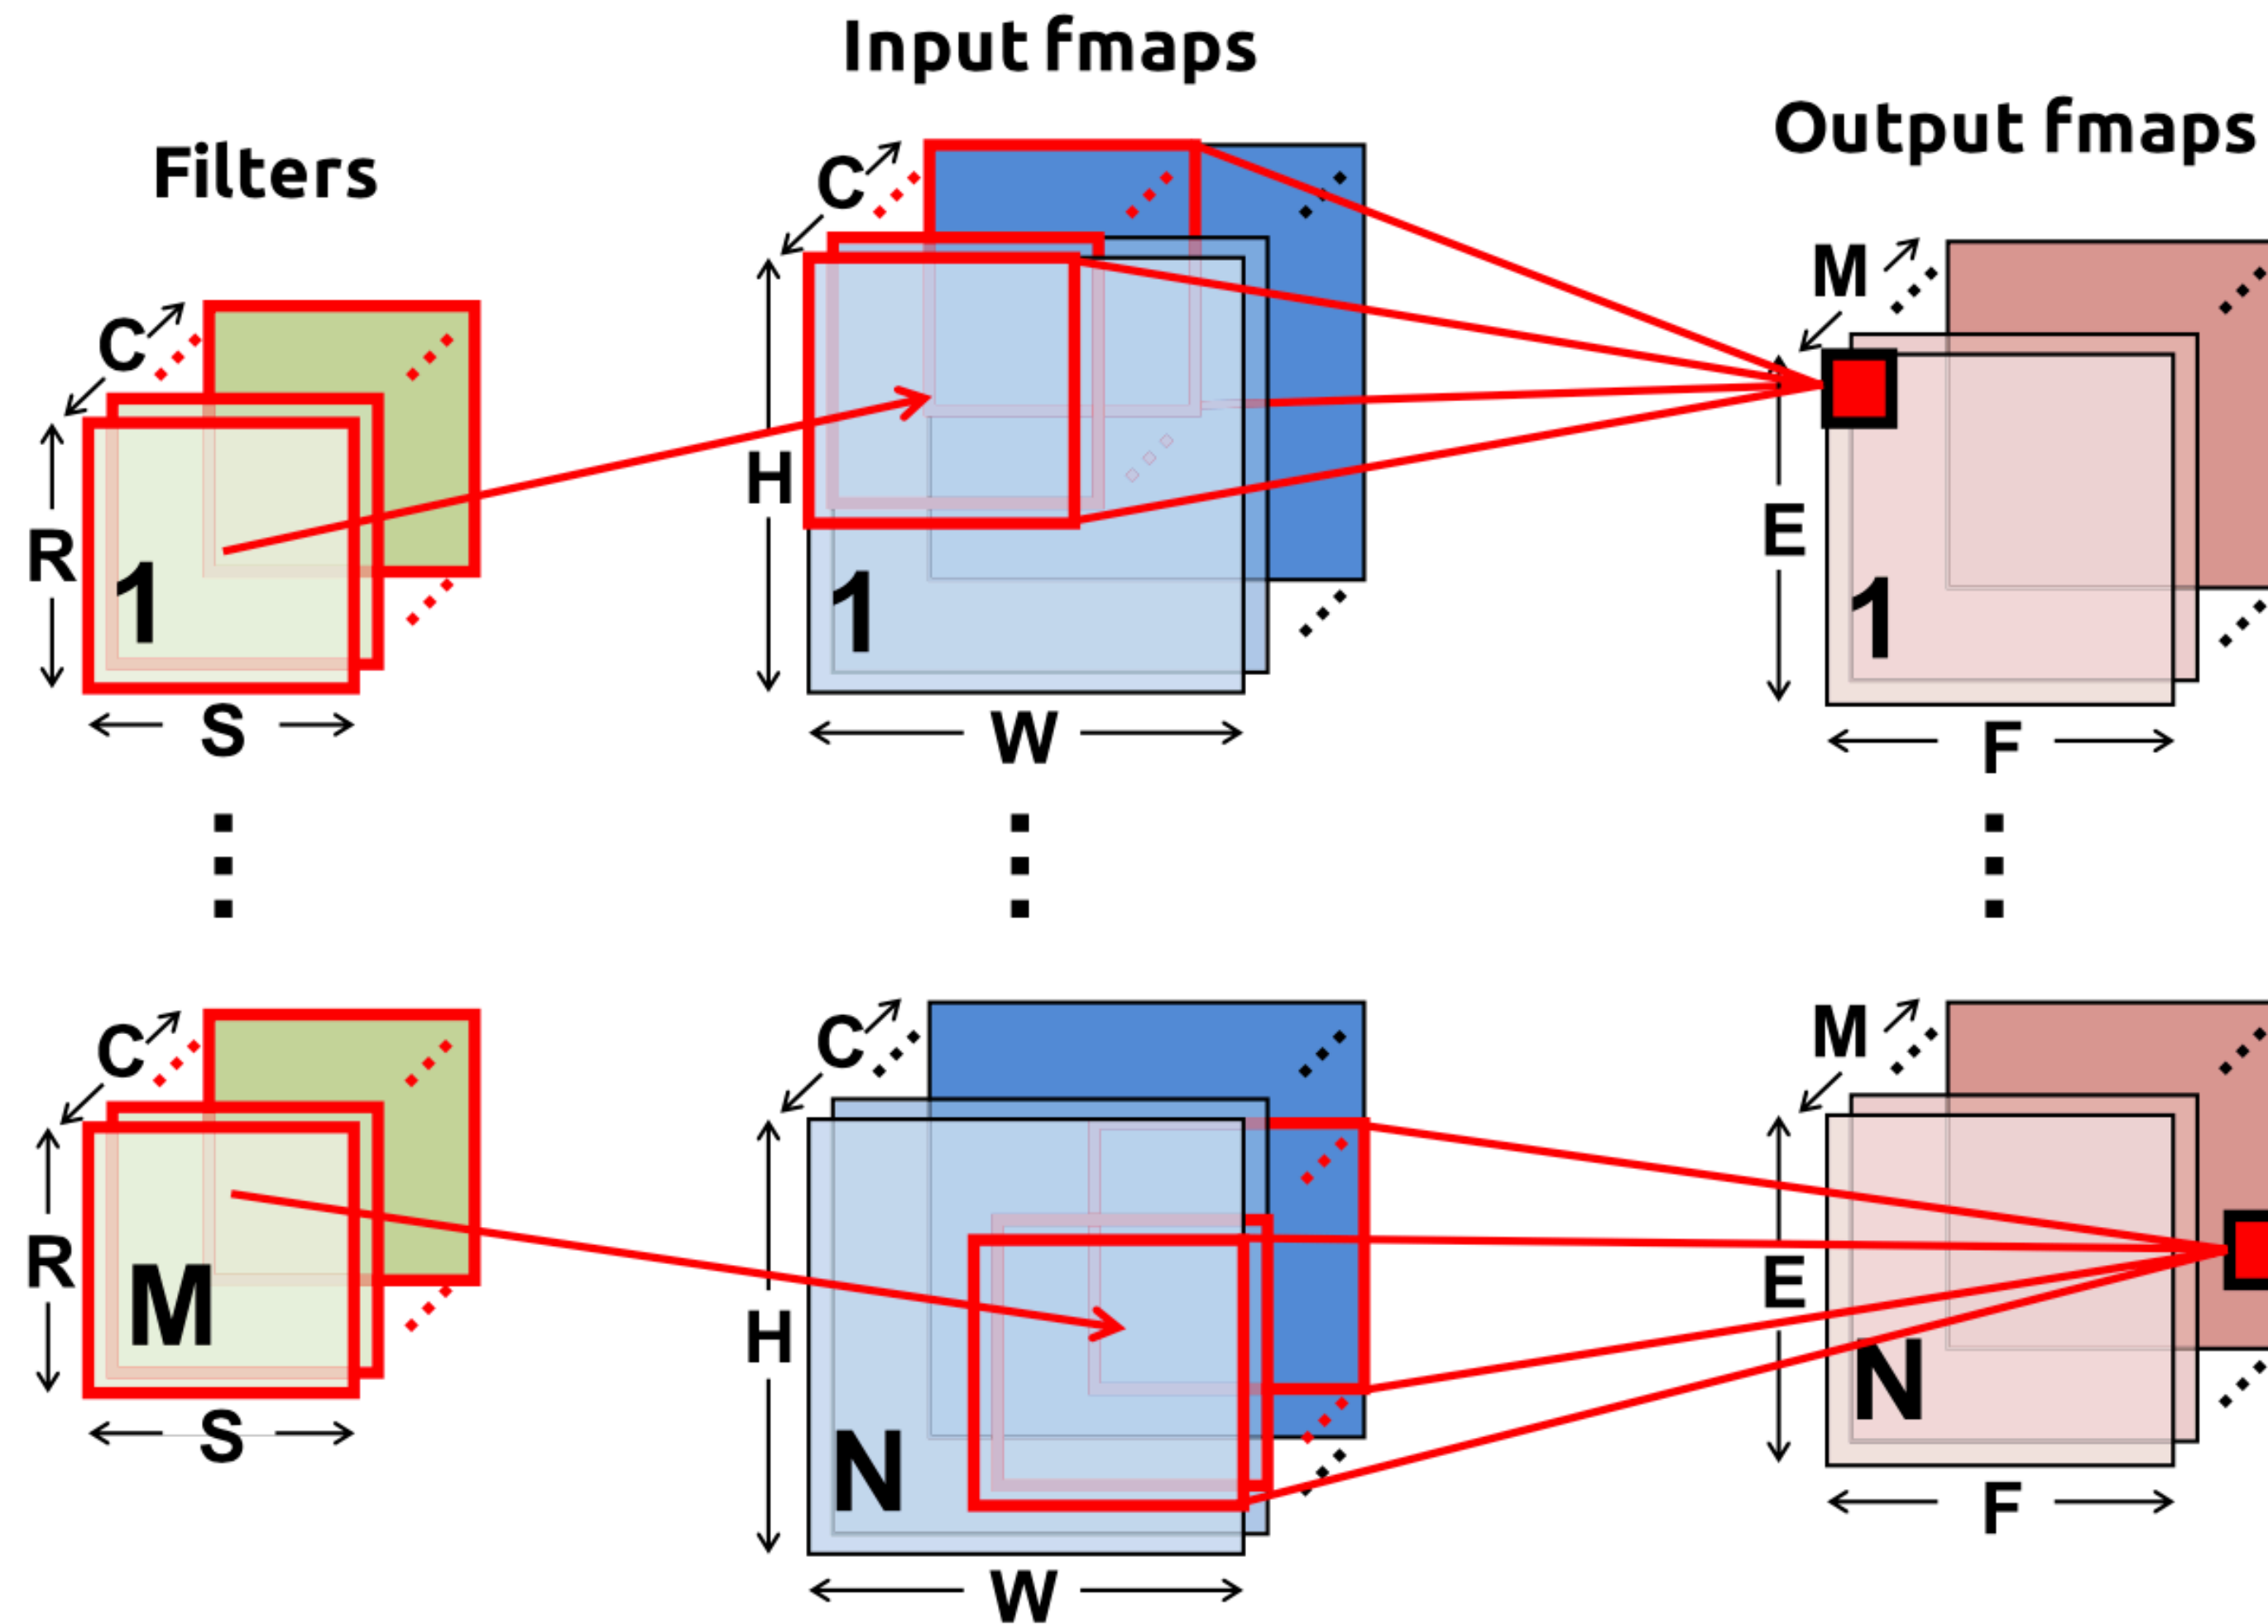
- Pose estimation

- Image captioning

Transfer learning: knowledge from one task is re-used for a different task

1. Pre-train your model on image classification tasks
2. Use it as a feature extractor for more complex tasks
3. Transfer learning then uses this model as base and adds more layers for more complex tasks; new layers will be fine tuned based on task-specific data

RECAP: CONVOLUTION



RECAP: CONVOLUTION

$$\mathbf{O}[z][u][x][y] = \sum_{k=0}^{C-1} \sum_{i=0}^{S-1} \sum_{j=0}^{R-1} \mathbf{I}[z][k][Ux + i][Uy + j] \cdot \mathbf{W}[u][k][i][j] + \mathbf{B}[u]$$

ofmap \mathbf{O} , ifmap \mathbf{I} , filters (weights) \mathbf{W} , and biases \mathbf{B}

ofmap = output filter map (output activations)

ifmap = input filter map (input activations)

$$E = (H - R + U)/U$$

$$F = (W - S + U)/U$$

N	Batch size (3D fmaps)	$0 \leq z \leq N$
M	number of 3D filters / number of ofmaps	$0 \leq u \leq M$
C	number of ifmap/filter channels	k
H / W	ifmap plane height/width	x resp y
R / S	filter plane height/width	x resp y
E / F	ofmap plane height/width	$0 \leq x \leq F, 0 \leq y \leq E$
U	stride	

RESOURCE-EFFICIENT MODEL ARCHITECTURES

Resource efficiency by model design

SQUEEZENET (2016)

Convolutional weights: $W_c = RSCM$

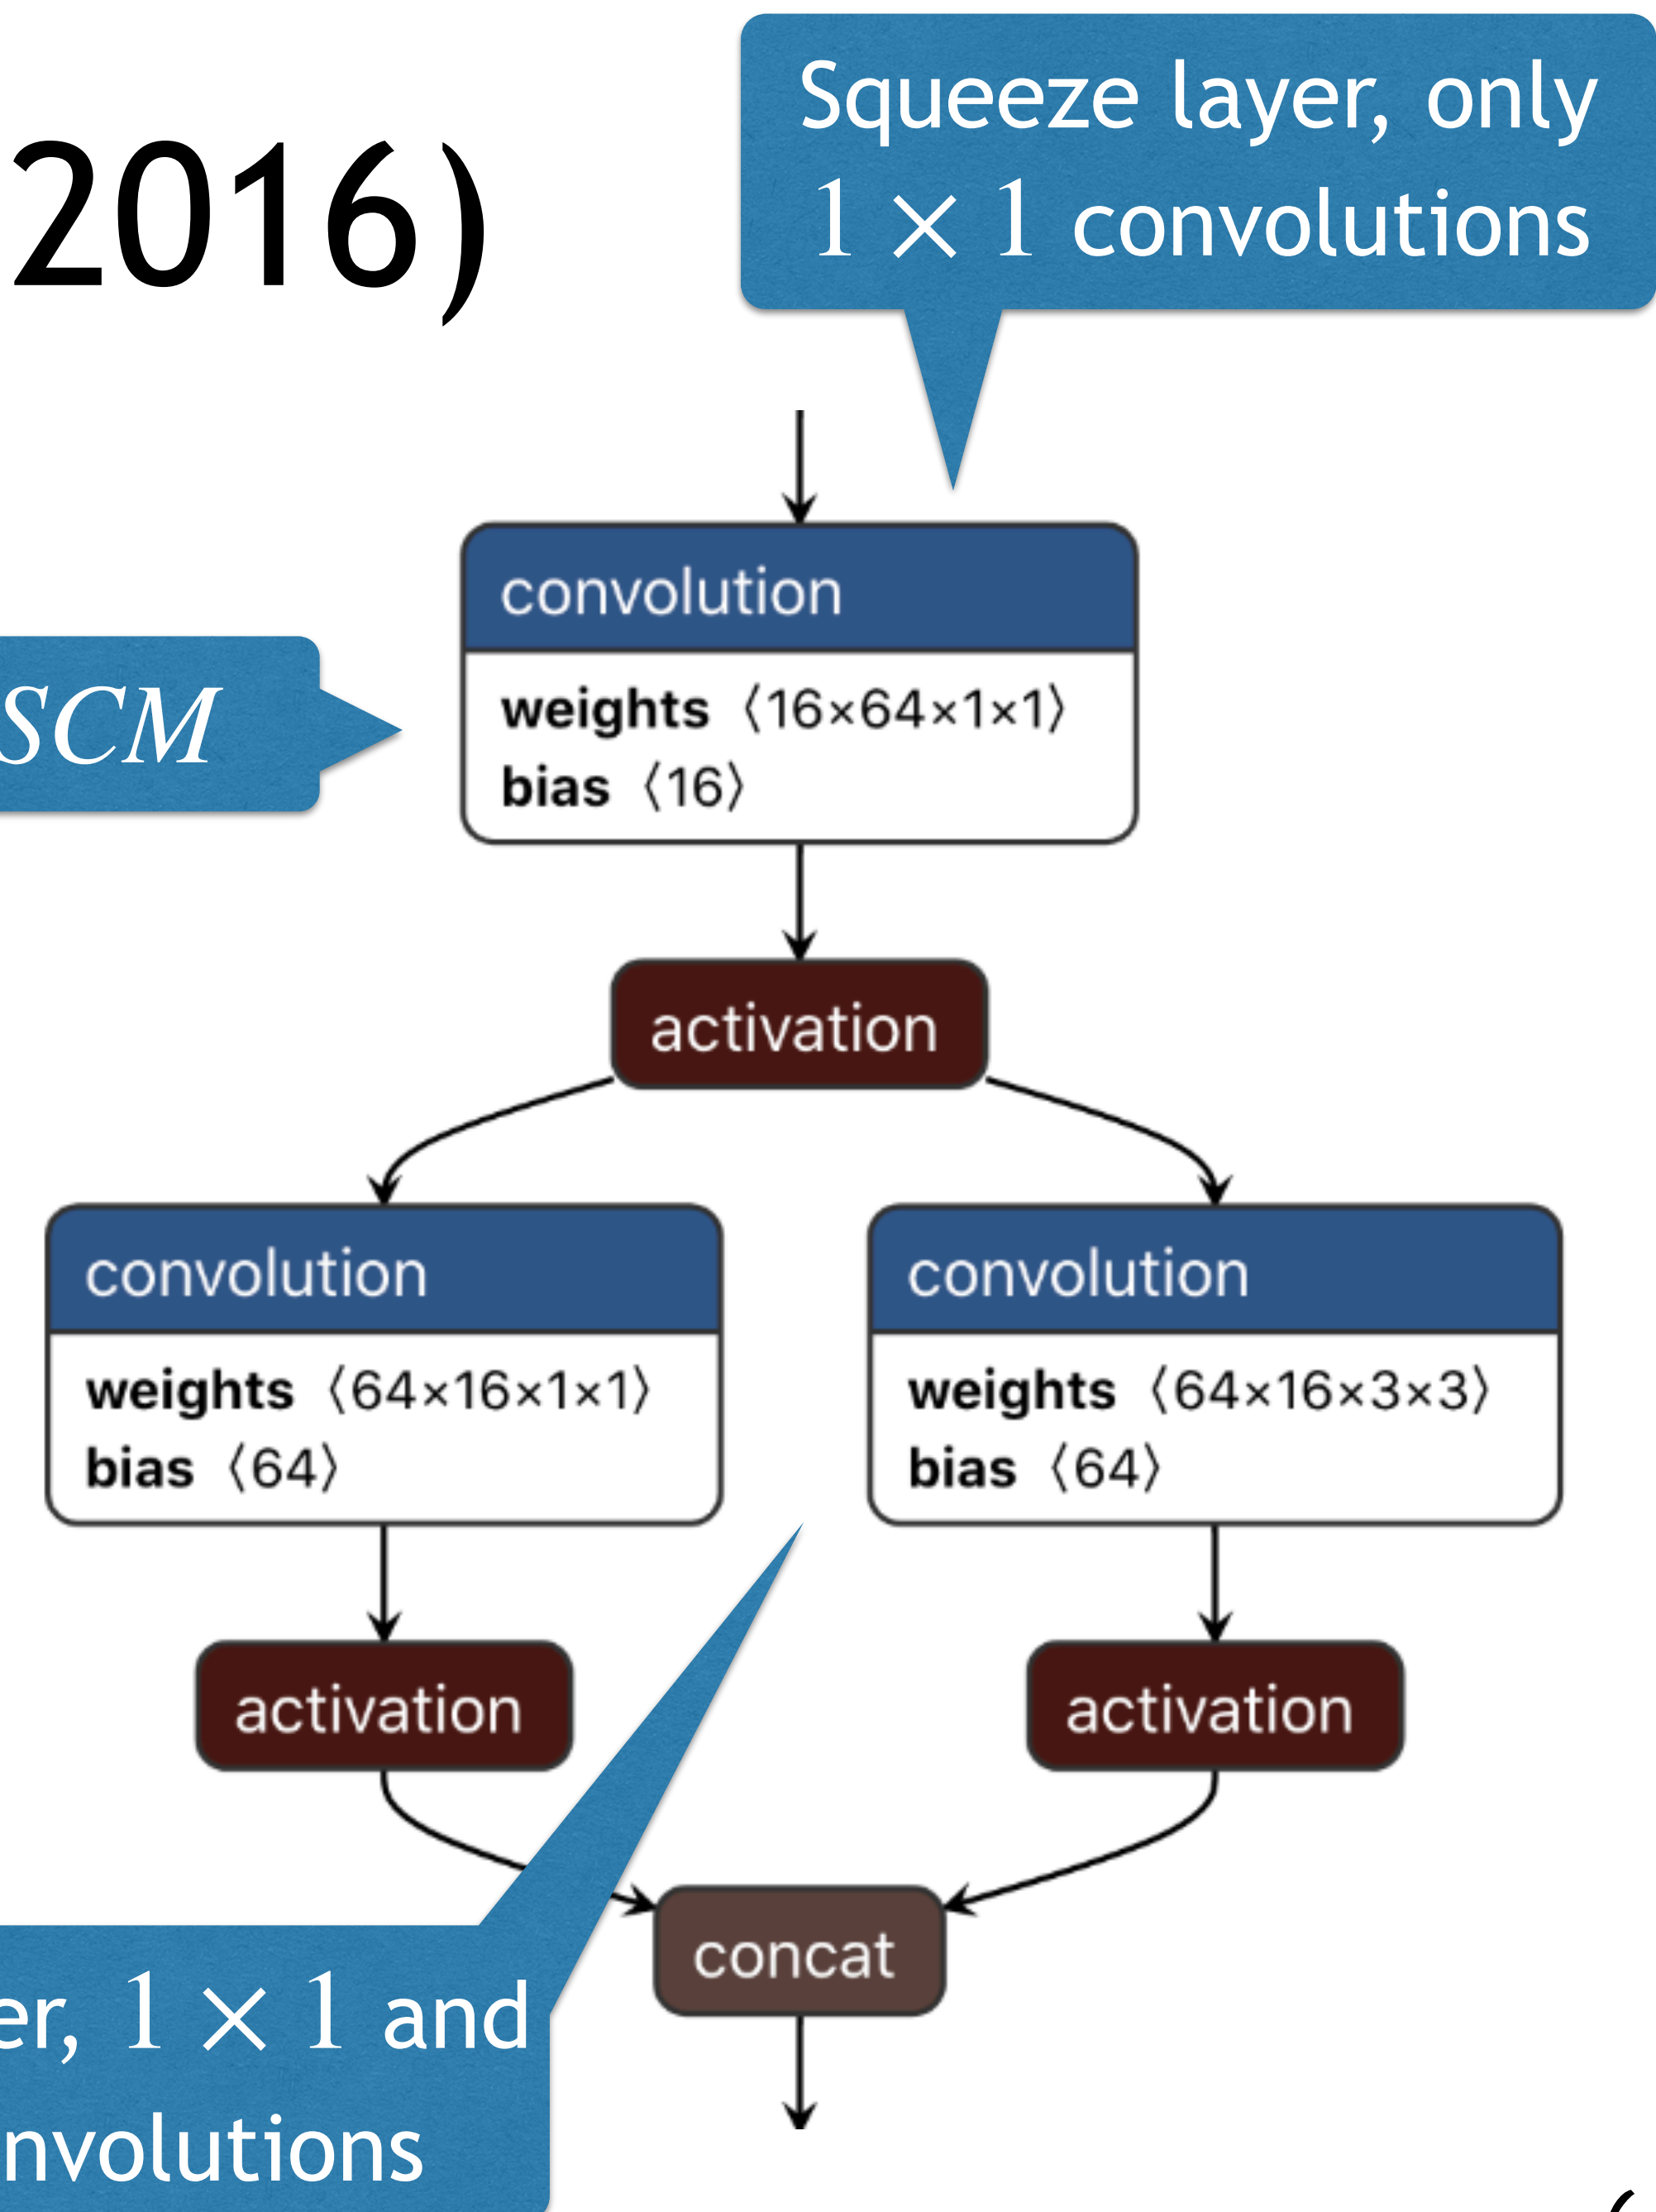
Three pillars

1. Replace 3x3 filters with 1x1 filters to reduce number of parameters (reduces R and S)
2. Decrease the number of input channels to 3x3 filters (reduces C)
3. Downsample late to maintain large activation maps (maintain accuracy)

Use 1. and 2. to implement a “fire module”

Example: $C=64 \rightarrow 16 \rightarrow 64$

Expand layer, 1×1 and 3×3 convolutions



SQUEEZENET ARCHITECTURE

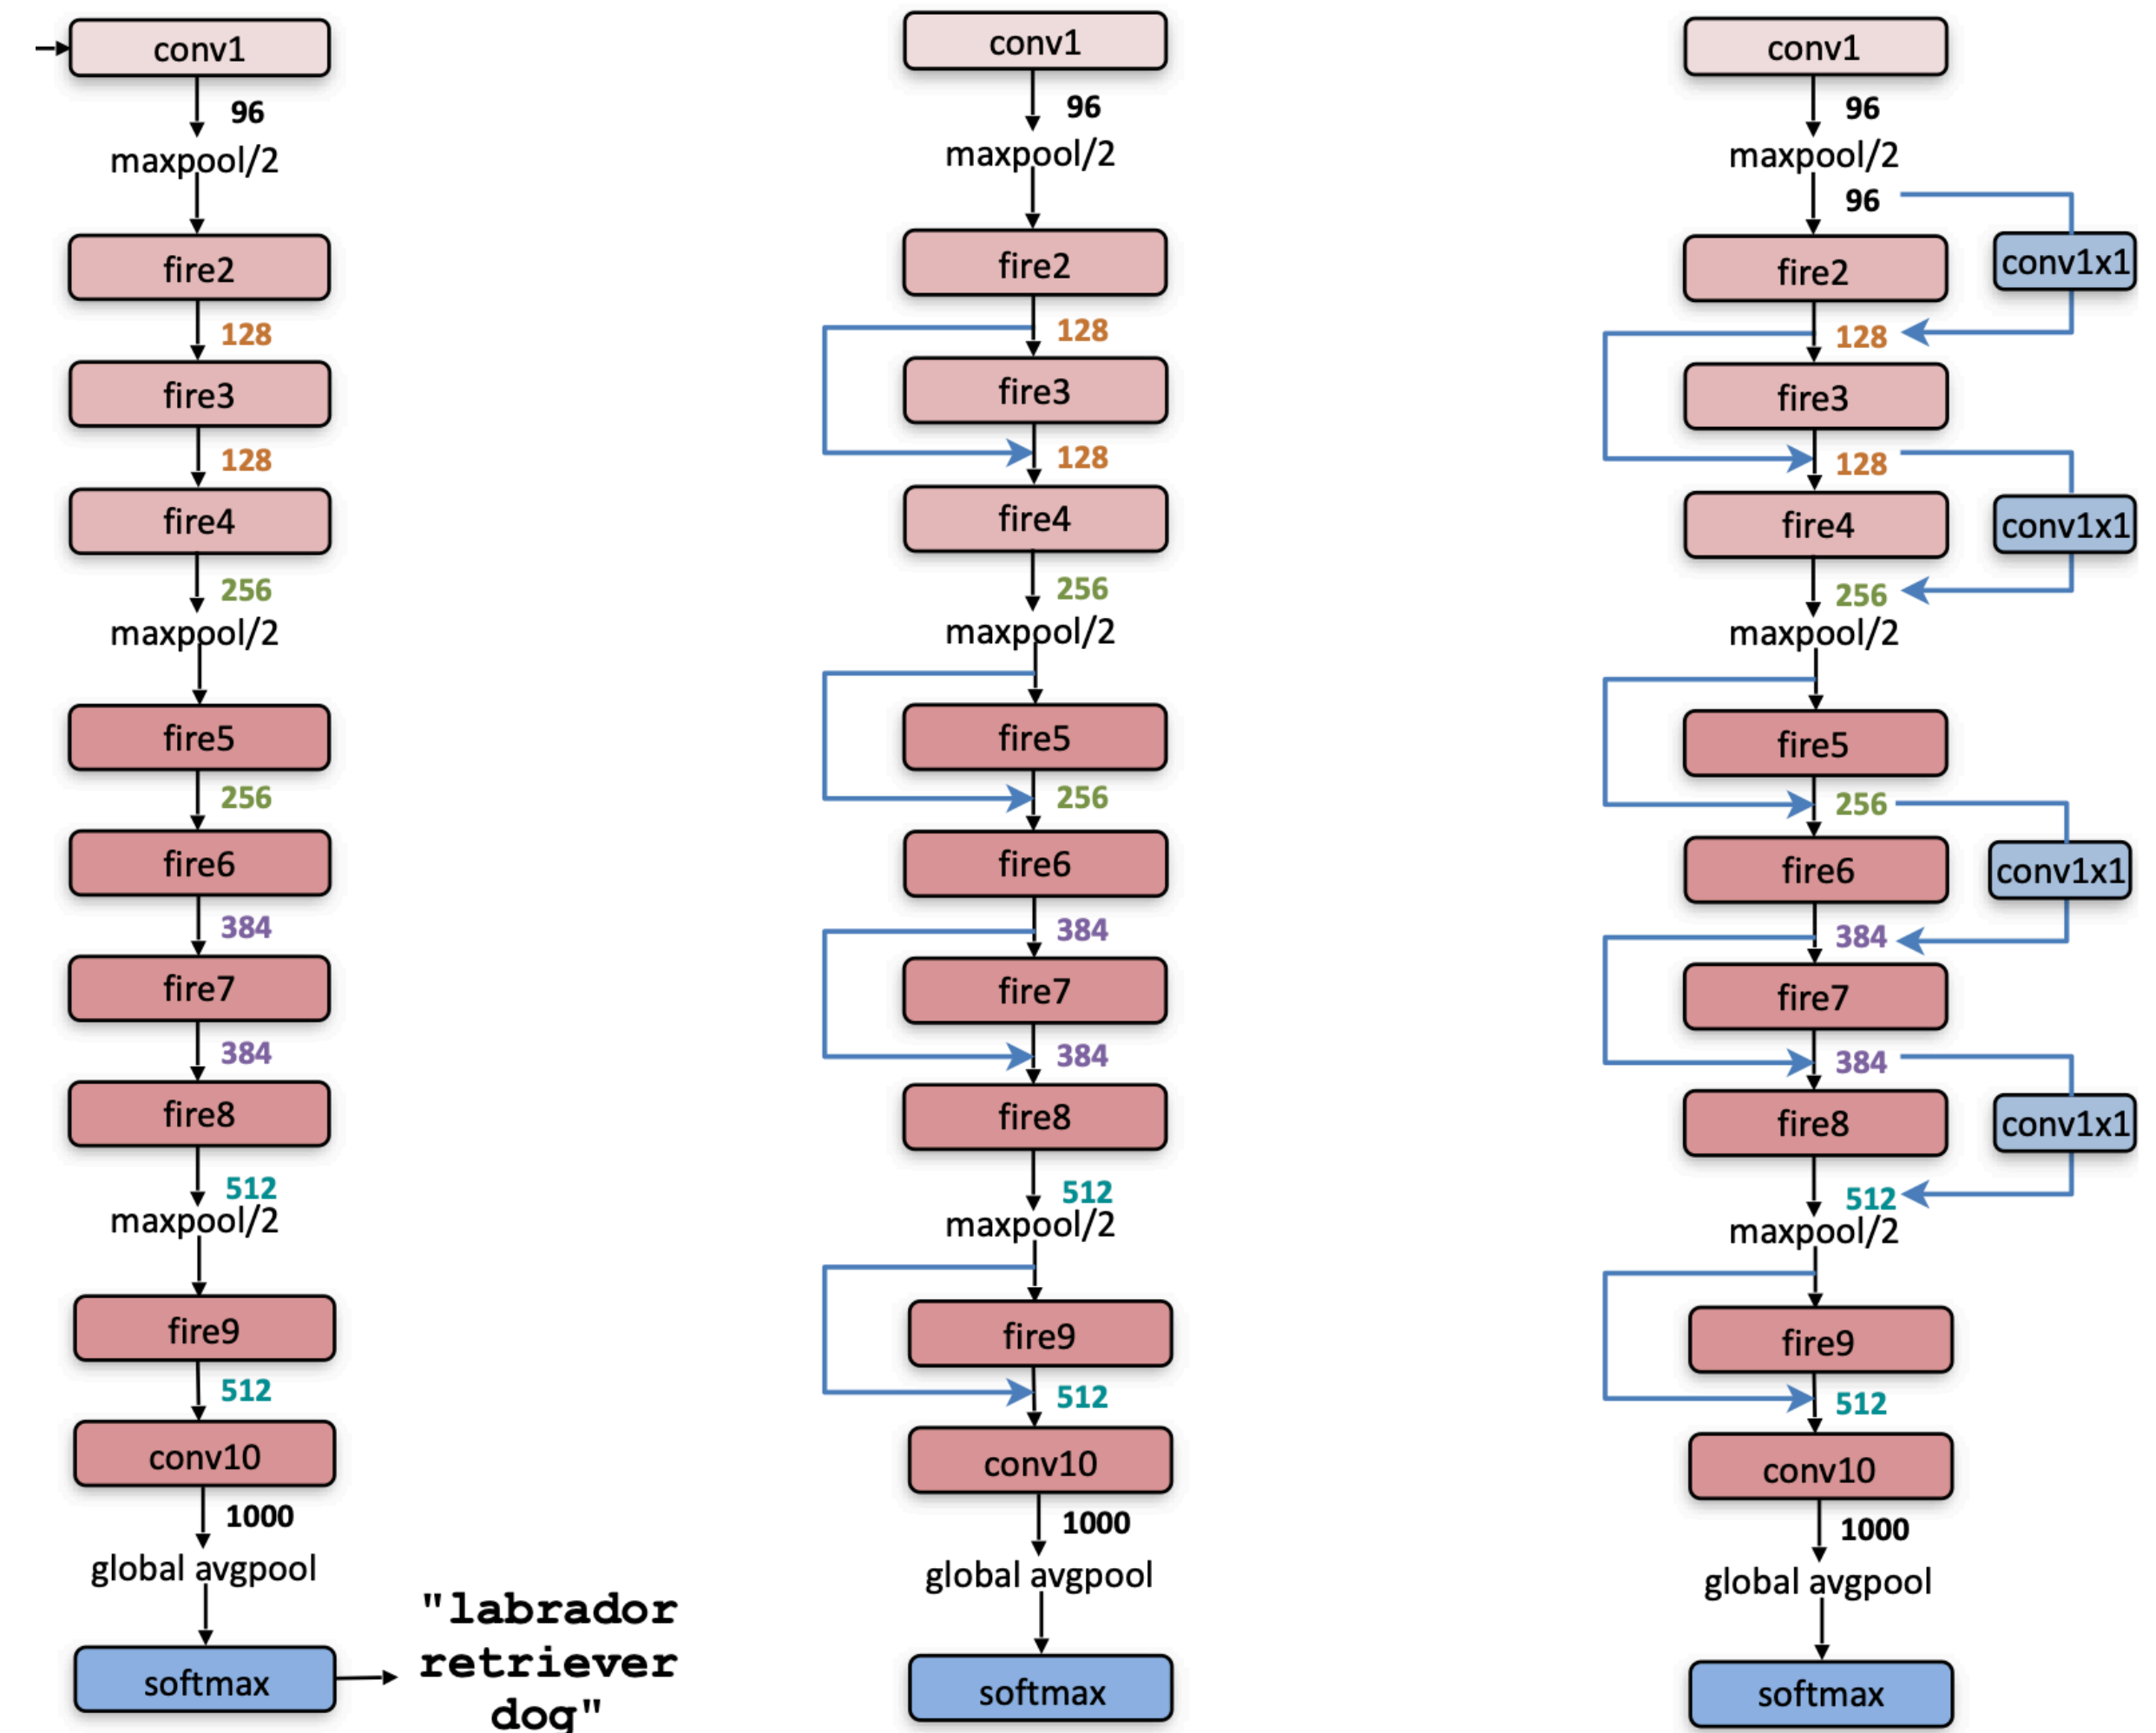
Different variants

No bypass, simple bypass, complex bypass

No FC layers

Max pooling (stride 2) relatively late

“Pooling layers reduce the size of the output map”



Forrest N. Iandola, Song Han, Matthew W. Moskewicz, Khalid Ashraf, William J. Dally, Kurt Keutzer, SqueezeNet: AlexNet-level accuracy with 50x fewer parameters and <0.5MB model size, <https://arxiv.org/abs/1602.07360>

SQUEEZENET RESULTS

CNN architecture	Compression Approach	Data Type	Original → Compressed Model Size	Reduction in Model Size vs. AlexNet	Top-1 ImageNet Accuracy	Top-5 ImageNet Accuracy
AlexNet	None (baseline)	32 bit	240MB	1x	57.2%	80.3%
AlexNet	SVD (Denton et al., 2014)	32 bit	240MB → 48MB	5x	56.0%	79.4%
AlexNet	Network Pruning (Han et al., 2015b)	32 bit	240MB → 27MB	9x	57.2%	80.3%
AlexNet	Deep Compression (Han et al., 2015a)	5-8 bit	240MB → 6.9MB	35x	57.2%	80.3%
SqueezeNet (ours)	None	32 bit	4.8MB	50x	57.5%	80.3%
SqueezeNet (ours)	Deep Compression	8 bit	4.8MB → 0.66MB	363x	57.5%	80.3%
SqueezeNet (ours)	Deep Compression	6 bit	4.8MB → 0.47MB	510x	57.5%	80.3%

AlexNet accuracy with 50x less parameters

~500x less parameters in combination with model compression

DeepCompression: pruning and quantization, but not inline with general-purpose processors

However: AlexNet performance is not state-of-the-art, SqueezeNet mainly important for its simplicity

Forrest N. Iandola, Song Han, Matthew W. Moskewicz, Khalid Ashraf, William J. Dally, Kurt Keutzer, SqueezeNet: AlexNet-level accuracy with 50x fewer parameters and <0.5MB model size, <https://arxiv.org/abs/1602.07360>

MOBILENET V1 (2017)

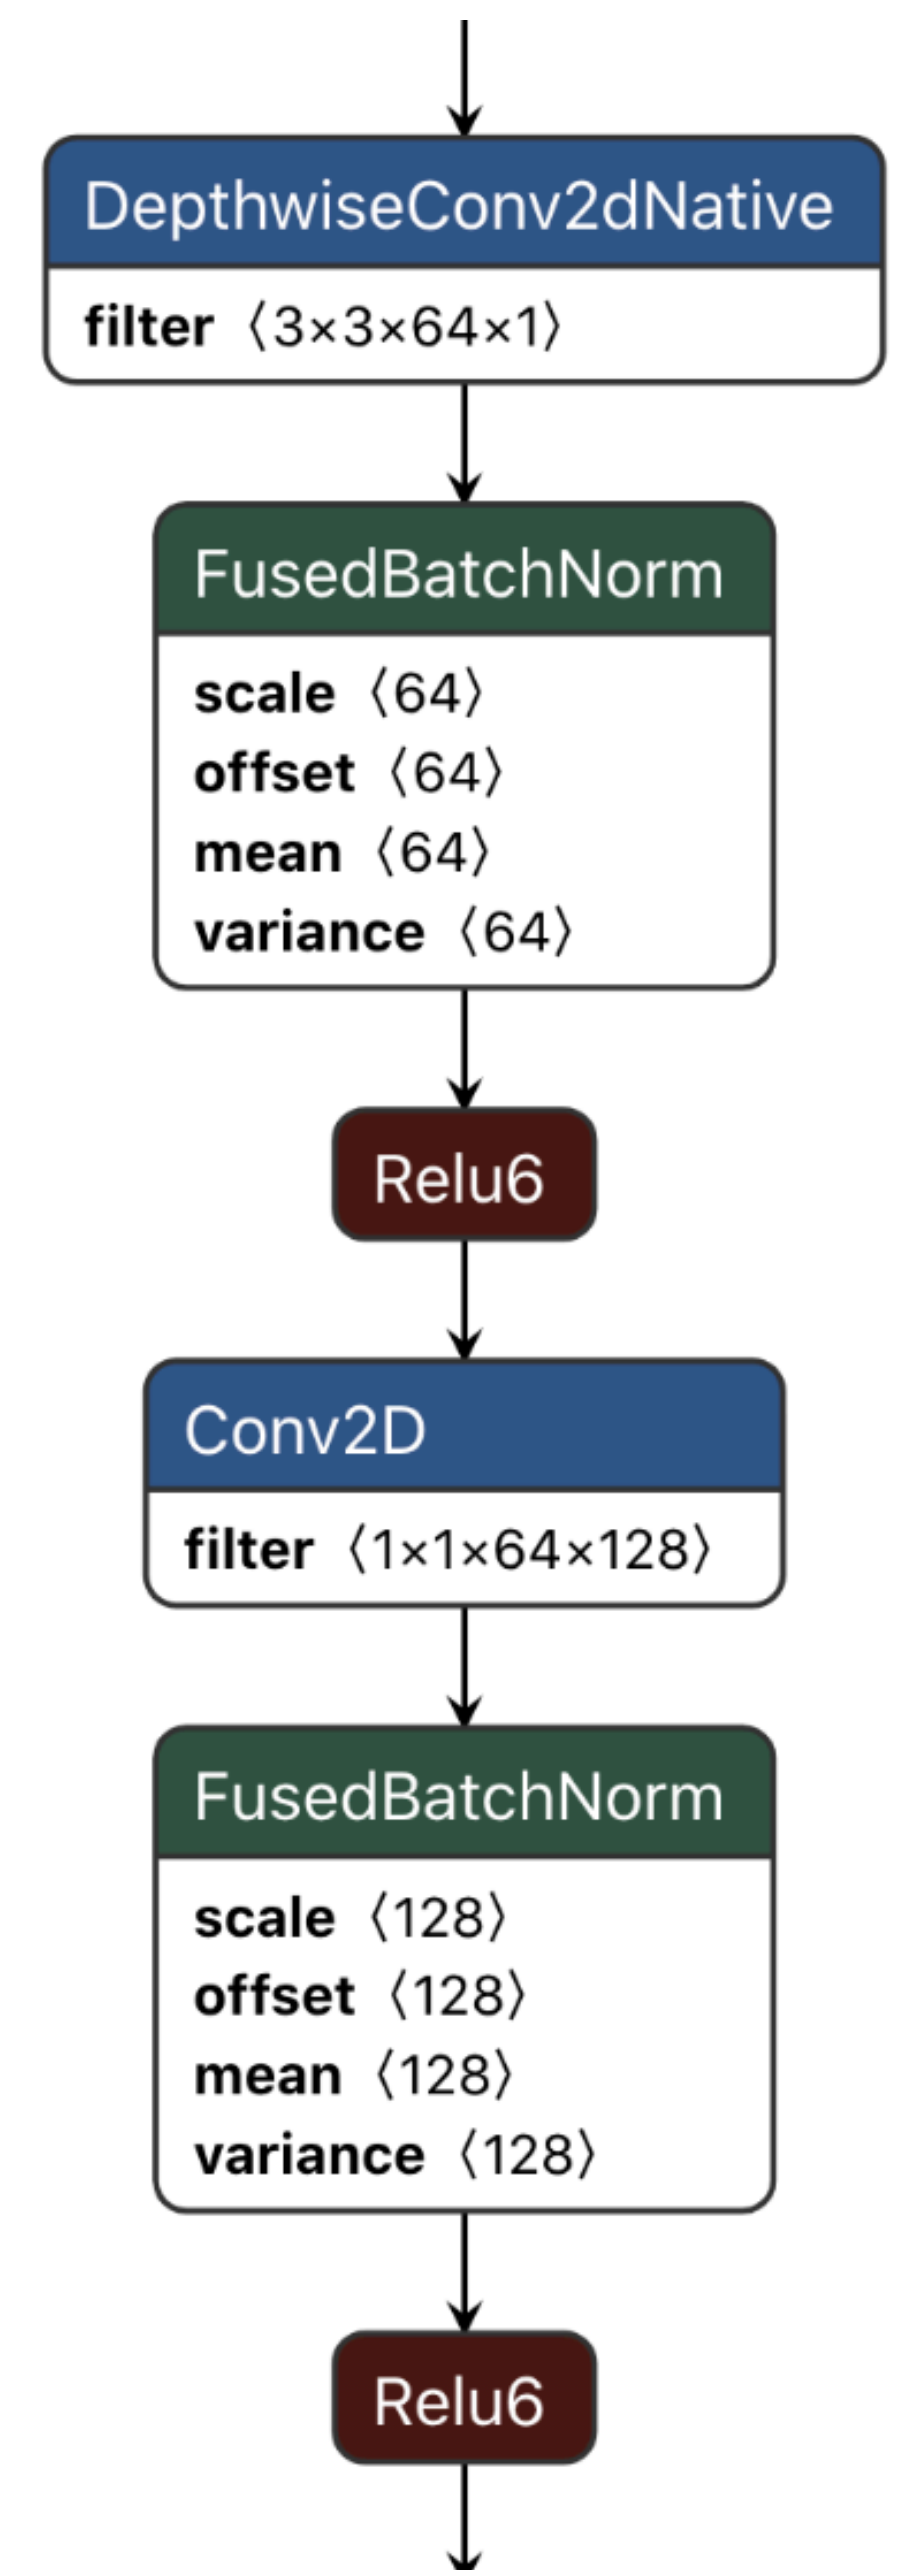
Core idea: replace expensive convolutions with cheaper alternatives

Depthwise separable convolution

Less trainable parameters

Less MACs

Separable convolutions?



(SPATIALLY) SEPARABLE CONVOLUTIONS

An matrix $K^{(M \times N)}$ is separable if it can be decomposed into $(M \times 1)$ and $(1 \times N)$ vectors

$$\begin{bmatrix} ax & ay & az \\ bx & by & bz \\ cx & cy & cz \end{bmatrix} = \begin{bmatrix} a \\ b \\ c \end{bmatrix} \cdot [x \ y \ z], \text{ e.g., } \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \cdot [1 \ 2 \ 1]$$

A convolution is then

$$\begin{aligned} o = \mathbf{I} * \mathbf{K} &= \begin{bmatrix} i_{1,1} & i_{1,2} & i_{1,3} \\ i_{2,1} & i_{2,2} & i_{2,3} \\ i_{3,1} & i_{3,2} & i_{3,3} \end{bmatrix} * \begin{bmatrix} ax & ay & az \\ bx & by & bz \\ cx & cy & cz \end{bmatrix} = \begin{bmatrix} i_{1,1} & i_{1,2} & i_{1,3} \\ i_{2,1} & i_{2,2} & i_{2,3} \\ i_{3,1} & i_{3,2} & i_{3,3} \end{bmatrix} * \left(\begin{bmatrix} a \\ b \\ c \end{bmatrix} \cdot [x \ y \ z] \right) \\ &= \left(\begin{bmatrix} i_{1,1} & i_{1,2} & i_{1,3} \\ i_{2,1} & i_{2,2} & i_{2,3} \\ i_{3,1} & i_{3,2} & i_{3,3} \end{bmatrix} * \begin{bmatrix} a \\ b \\ c \end{bmatrix} \right) * [x \ y \ z] \end{aligned}$$

* is a convolution operator,
not a multiplication

Reduces computational complexity from MN down to $M + N$ multiplications

However, consider increase in space complexity - why?

More intermediate results

(SPATIALLY) SEPARABLE CONVOLUTIONS - EXAMPLE

$$\begin{aligned} \mathbf{I} * \mathbf{K} &= \begin{bmatrix} 50 & 60 & 70 \\ 80 & 90 & 100 \\ 110 & 120 & 130 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \\ &= \begin{bmatrix} 50 & 60 & 70 \\ 80 & 90 & 100 \\ 110 & 120 & 130 \end{bmatrix} * \left(\begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \cdot [1 \ 2 \ 1] \right) \\ &= \left(\begin{bmatrix} 50 & 60 & 70 \\ 80 & 90 & 100 \\ 110 & 120 & 130 \end{bmatrix} * \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \right) * [1 \ 2 \ 1] \end{aligned}$$

* is a convolution operator,
not a multiplication

$$\begin{aligned} &= [50 \cdot 1 + 80 \cdot 2 + 110 \cdot 1 \quad 60 \cdot 1 + 90 \cdot 2 + 120 \cdot 1 \quad 70 \cdot 1 + 100 \cdot 2 + 130 \cdot 1] \\ &\quad * [1 \ 2 \ 1] \\ &= [320 \ 360 \ 400] * [1 \ 2 \ 1] = [320 \cdot 1 + 360 \cdot 2 + 400 \cdot 1] = 1440 \end{aligned}$$

DEPTHWISE SEPARABLE CONVOLUTIONS

Spatially separable convolutions are not used often in ML as kernel search is limited

Instead, consider $M R \times S \times C$ convolutions

E.g., transform a $7 \times 7 \times 3$ input shape into a $5 \times 5 \times 128$ output shape

Using 128 kernels of size $3 \times 3 \times 3$

Step 1: replace a $R \times S \times C$ convolution with $C R \times S \times 1$ (separable) convolutions

$3 \times 3 \times 3 \rightarrow 3$ kernels of size $3 \times 3 \times 1$

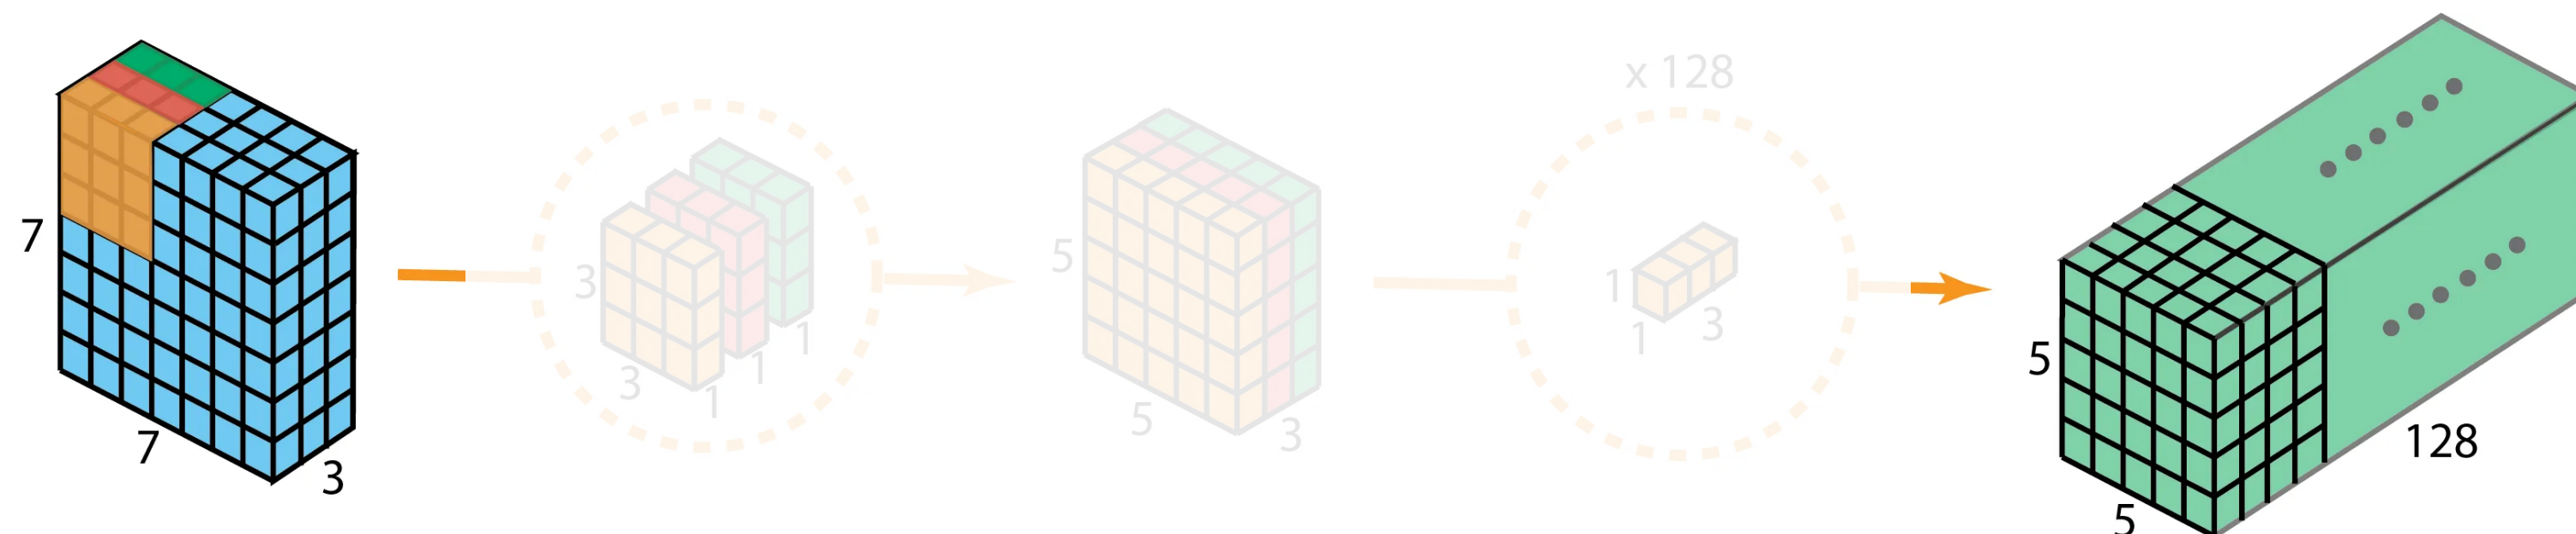
Output shape is identical by stacking output fmaps

pointwise
convolution

Step 2: as cross-channel correlation is lost, add afterwards $M 1 \times 1 \times C$ convolutions

One convolution then transforming the $7 \times 7 \times 3$ shape into a $7 \times 7 \times 1$ shape

All M convolutions then result in a $7 \times 7 \times 128$ shape



DEPTHWISE SEPARABLE CONVOLUTIONS - EFFICIENCY

Plain convolution: $MAC_c = (EF \cdot RSC) \cdot M$

For $E = (H - R + U)/U$ and $F = (W - S + U)/U$

For the following, assume a unit stride ($U = 1$) and zero padding

Then: $MAC_c = ((H - R + 1)(W - S + 1) \cdot RSC) \cdot M$

Step 1: replace $M R \times S \times C$ convolutions with $C R \times S \times 1$ convolutions

$$MAC_{cds1} = ((H - R + 1)(W - S + 1) \cdot RS \cdot 1) \cdot C$$

Step 2: followed by $M 1 \times 1 \times C$ convolutions

$$MAC_{cds1} = ((H - R + 1)(W - S + 1) \cdot 1 \cdot 1 \cdot C) \cdot M$$

$$MAC_{cds} = MAC_{cds1} + MAC_{cds2} = (H - R + 1)(W - S + 1) \cdot (RSC + CM)$$

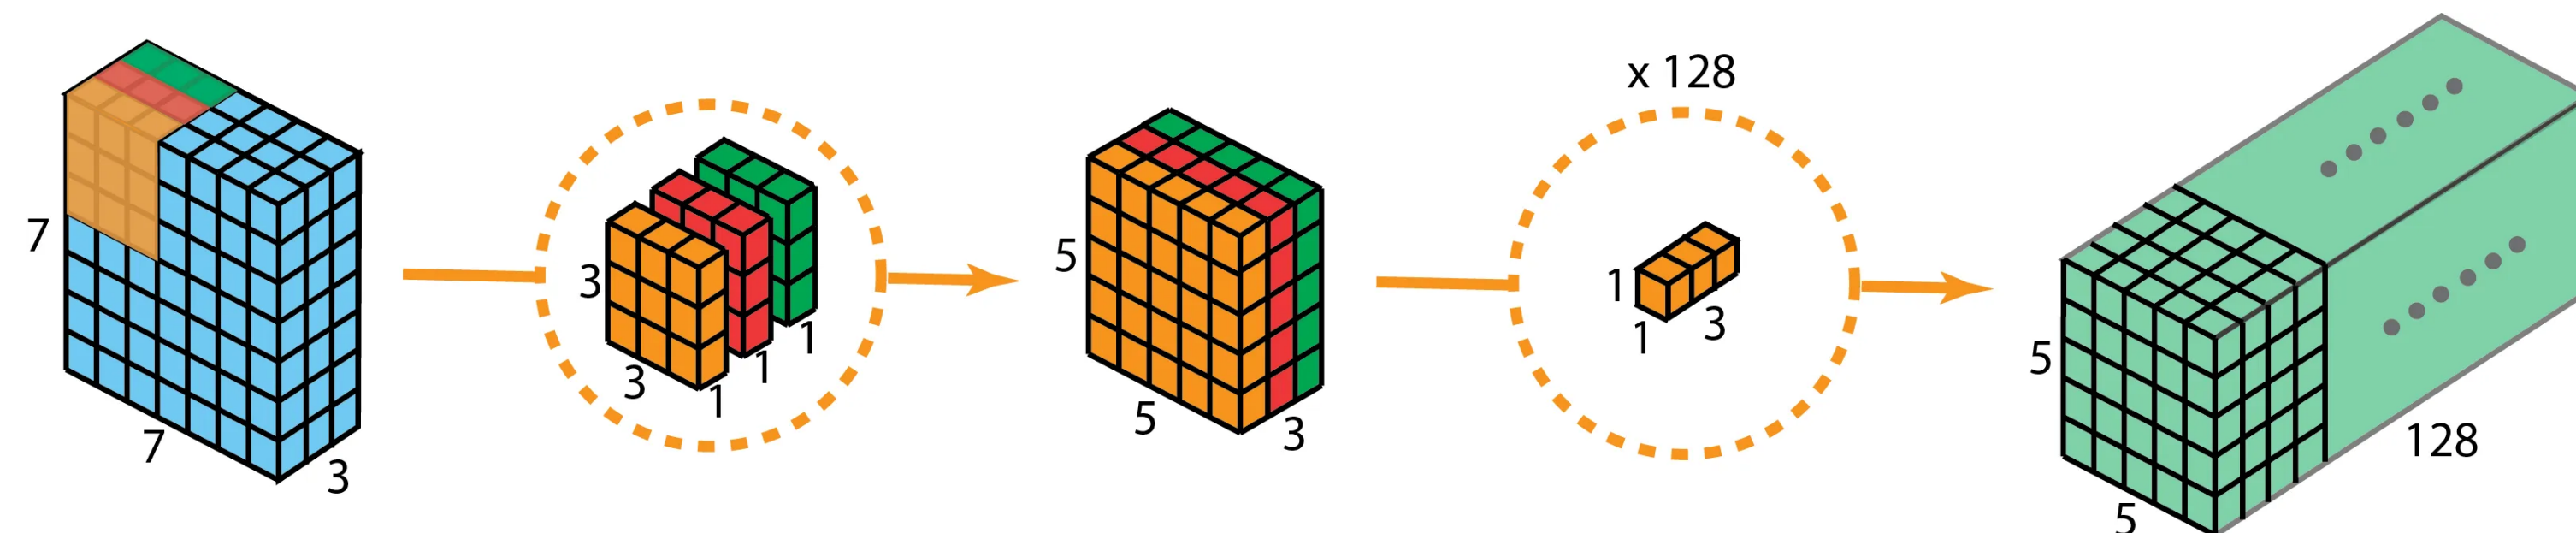
DEPTHWISE SEPARABLE CONVOLUTIONS - EFFICIENCY EXAMPLE

$$MAC_c = ((H - R + 1)(W - S + 1) \cdot RSC) \cdot M = 86400$$

$$MAC_{c ds} = ((H - R + 1)(W - S + 1) \cdot (RSC + CM)) = 10275$$

$$r = \frac{MAC_c}{MAC_{c ds}} = \frac{RSM}{RSC + CM} = \frac{1}{M} + \frac{1}{RS}$$

Modern architectures: $M \gg R, S \Rightarrow r \approx \frac{1}{RS}$



MOBILENET V1 (2017)

Core idea: replace expensive convolutions with cheaper alternatives

Depthwise separable convolution

Less trainable parameters

Less MACs

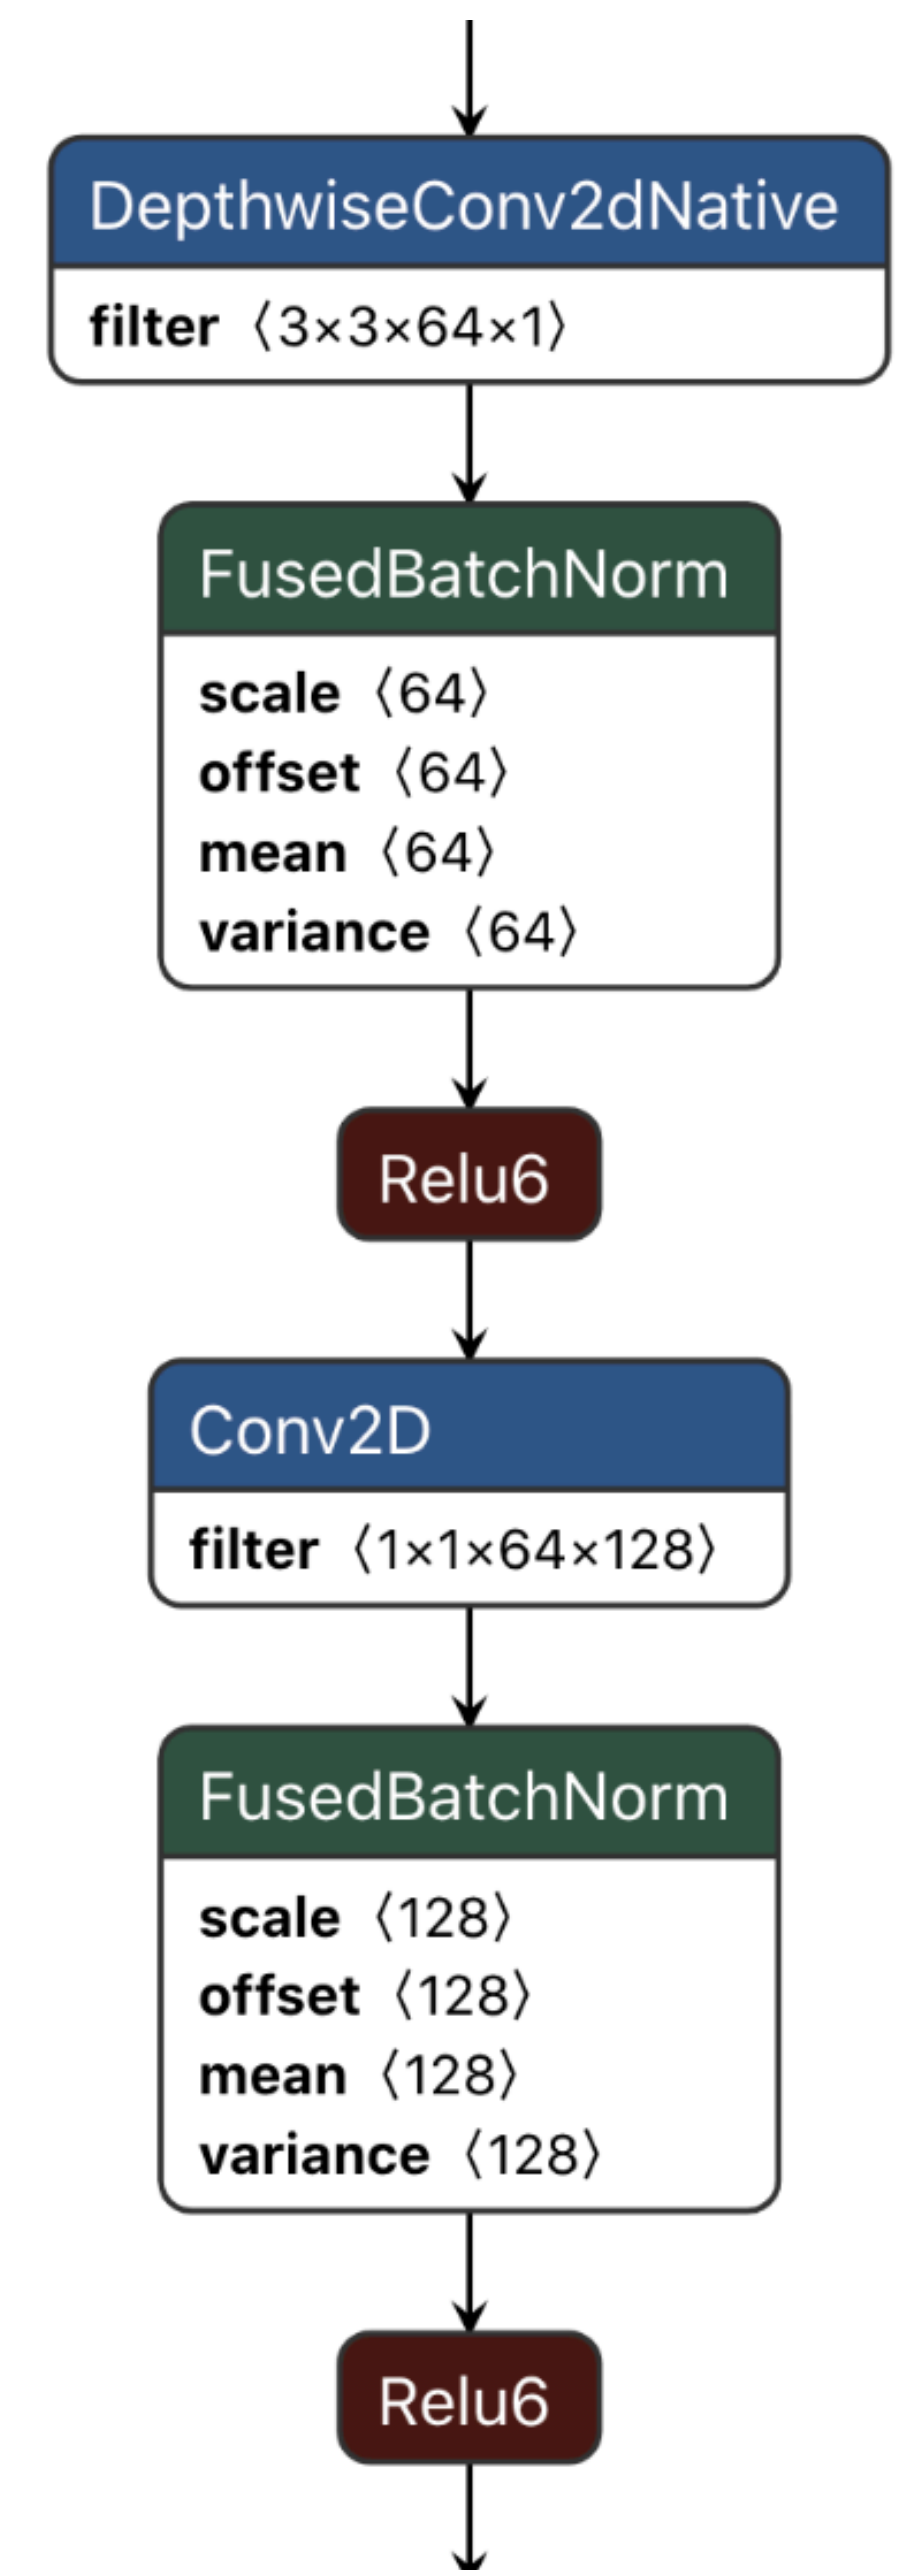
Depthwise separable convolutions

Save: MACs, save parameters

Cost: parameters (less model capacity), activations

Width multiplier $\alpha \in (0,1]$ to control the input channel depth

Thinner models based on typical settings of 1, 0.75, 0.5, and 0.25



MOBILENET V1 (2017)

Table 4. Depthwise Separable vs Full Convolution MobileNet

Model	ImageNet	Million	Million
	Accuracy	Mult-Adds	Parameters
Conv MobileNet	71.7%	4866	29.3
MobileNet	70.6%	569	4.2

Table 6. MobileNet Width Multiplier

Width Multiplier	ImageNet	Million	Million
	Accuracy	Mult-Adds	Parameters
1.0 MobileNet-224	70.6%	569	4.2
0.75 MobileNet-224	68.4%	325	2.6
0.5 MobileNet-224	63.7%	149	1.3
0.25 MobileNet-224	50.6%	41	0.5

Outperforms various model alternatives based on accuracy, MACs, and parameter count

GoogLeNet, VGG16, AlexNet, SqueezeNet ($\alpha = 0.5$), Inception v3 (“Stanford dogs” dataset)

MobileNet-SSD for object detection based on MS-COCO

MOBILENET V2 (2018)

Rearranged architecture based on building blocks

1. Expansion layer based on 1×1 convolution to increase channel depth
2. Depthwise convolutional layer
3. Bottleneck layer to reduce channel depth

Skip connection if input/output dimensions match

“Inverted residuals” as ResNet used skip connections based on many channels (reducing width inside a block)

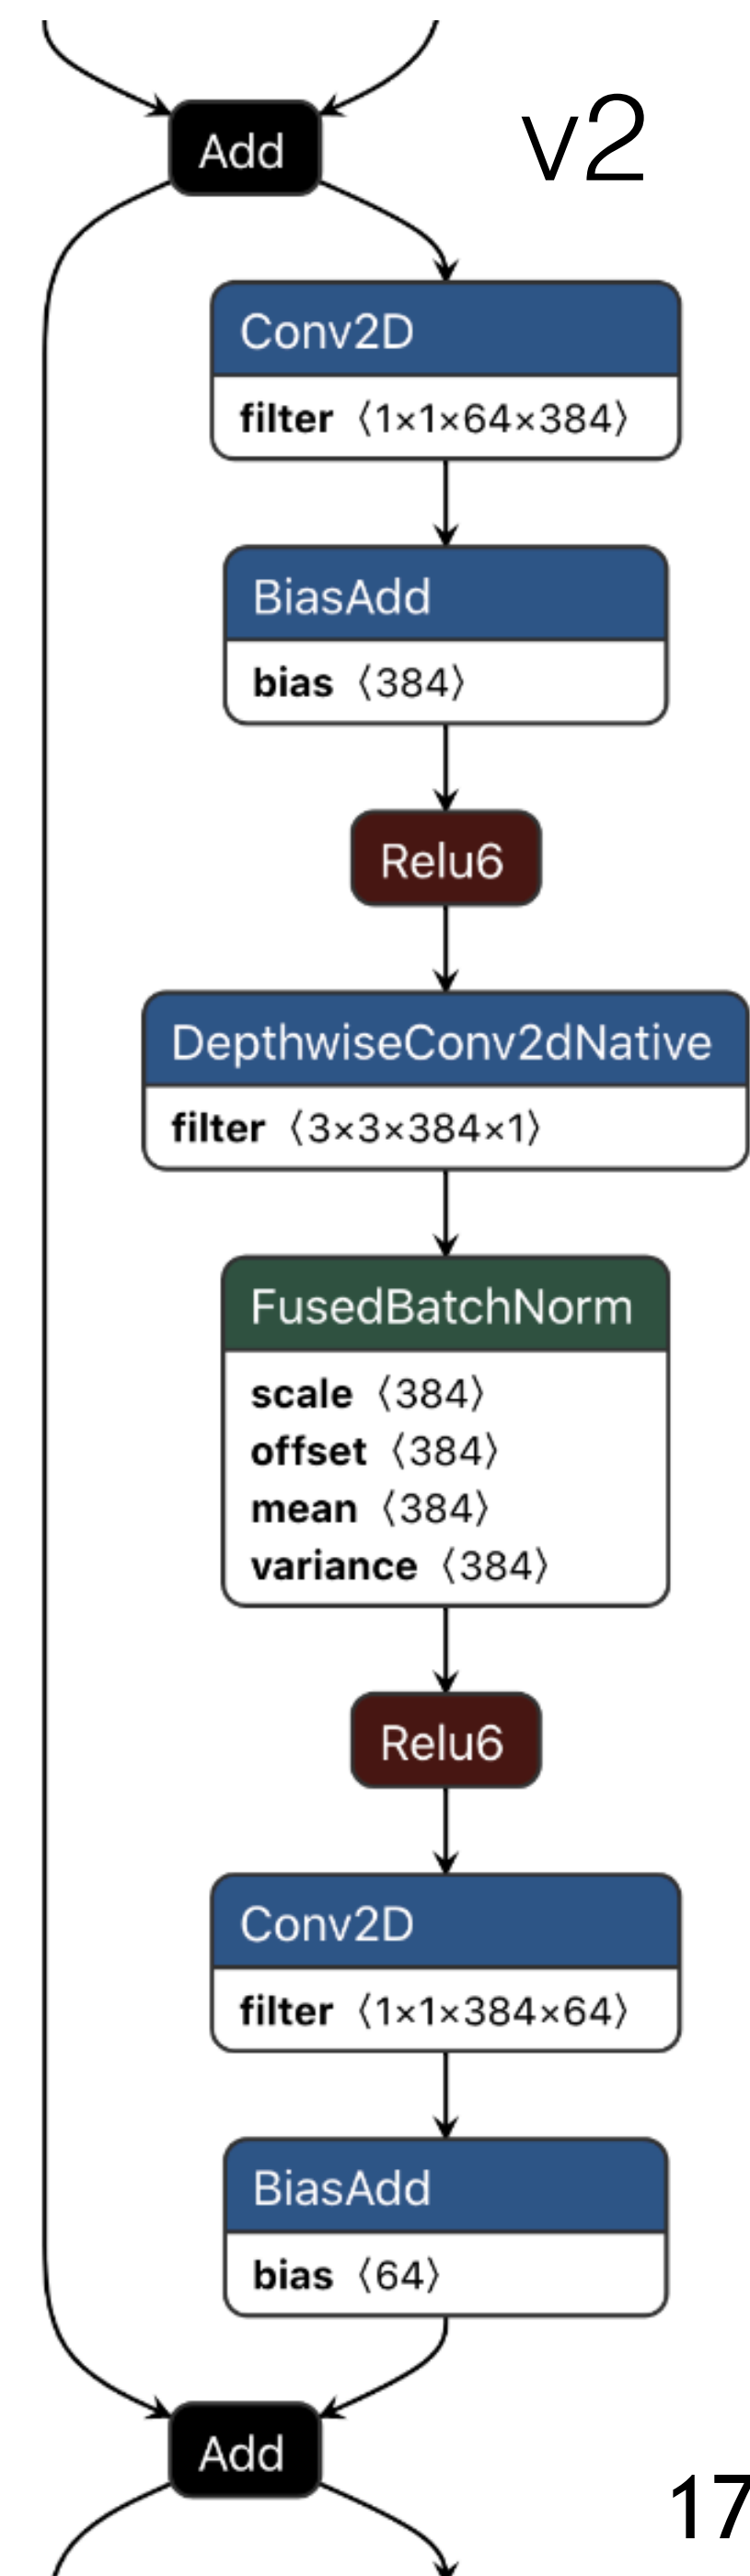
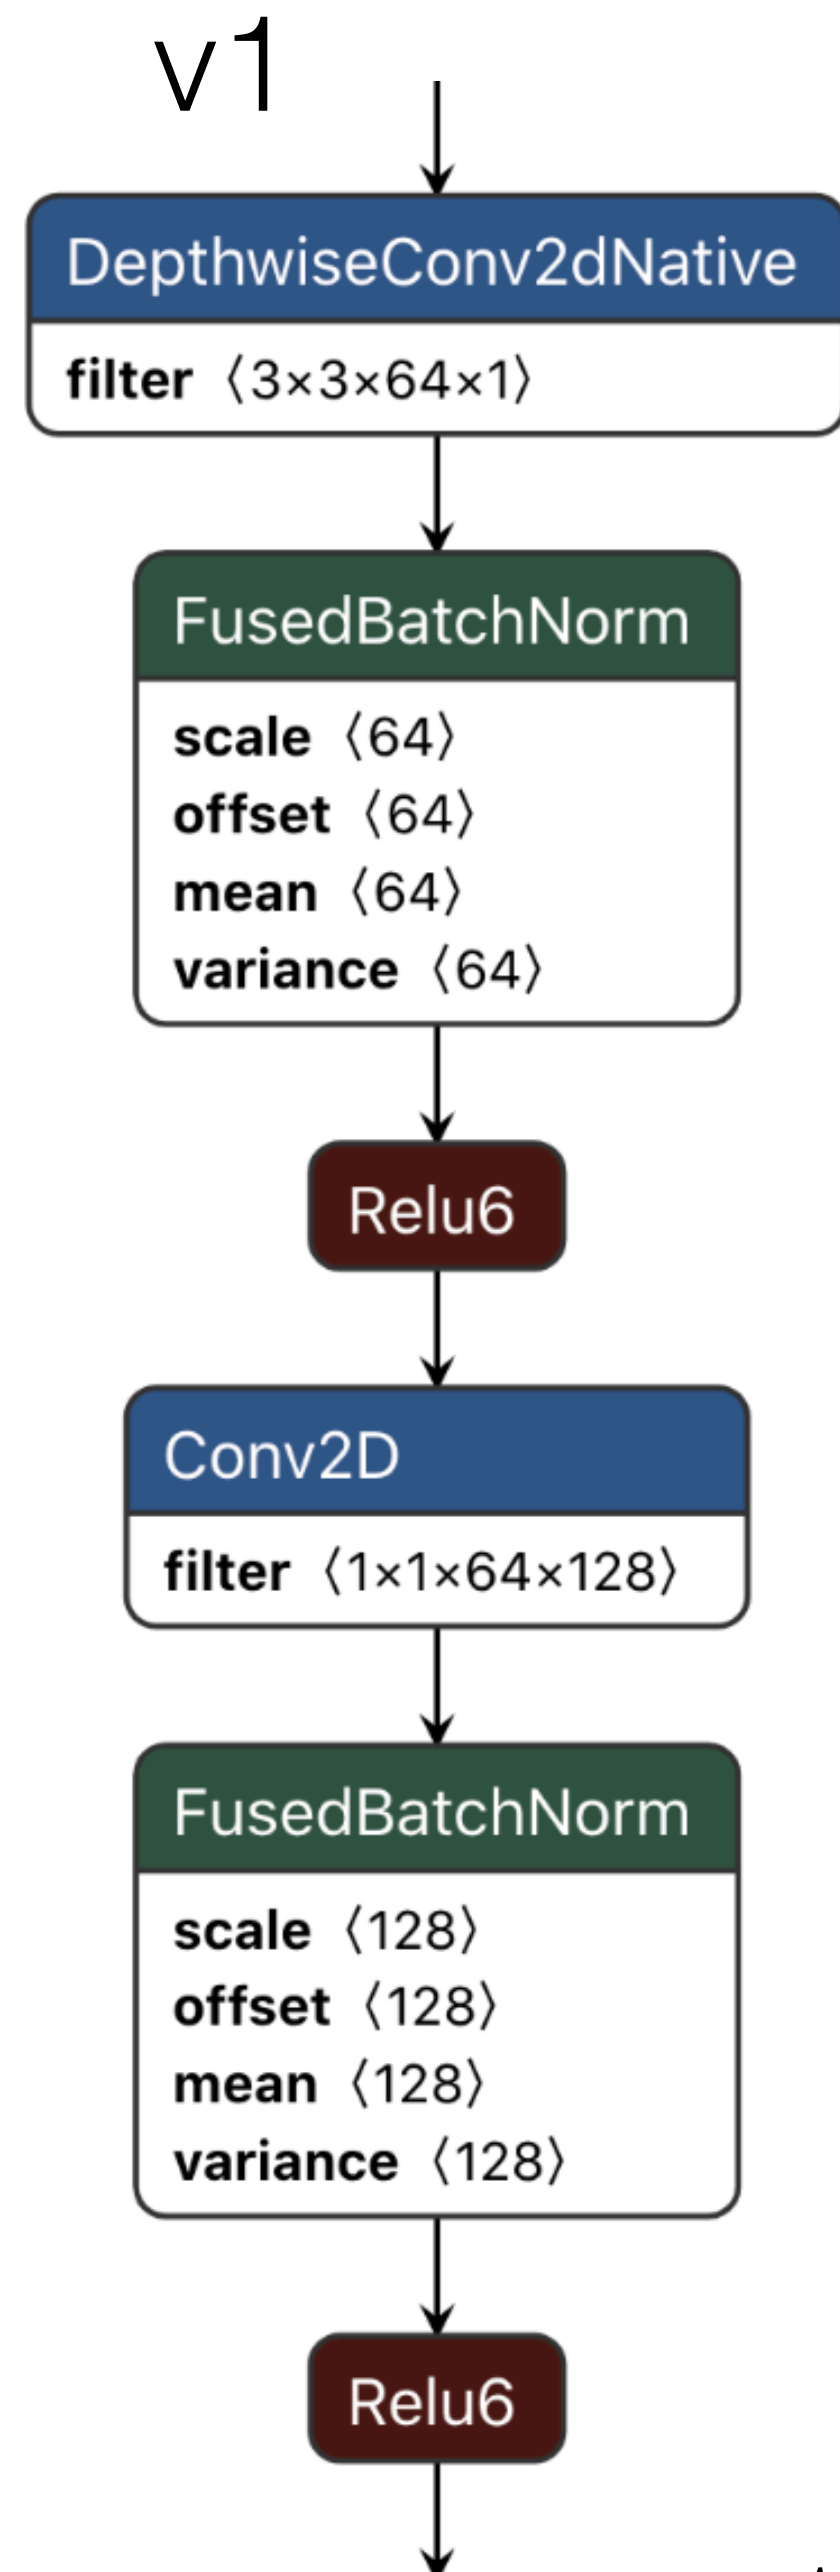
More memory efficient

No nonlinearity behind bottleneck layer

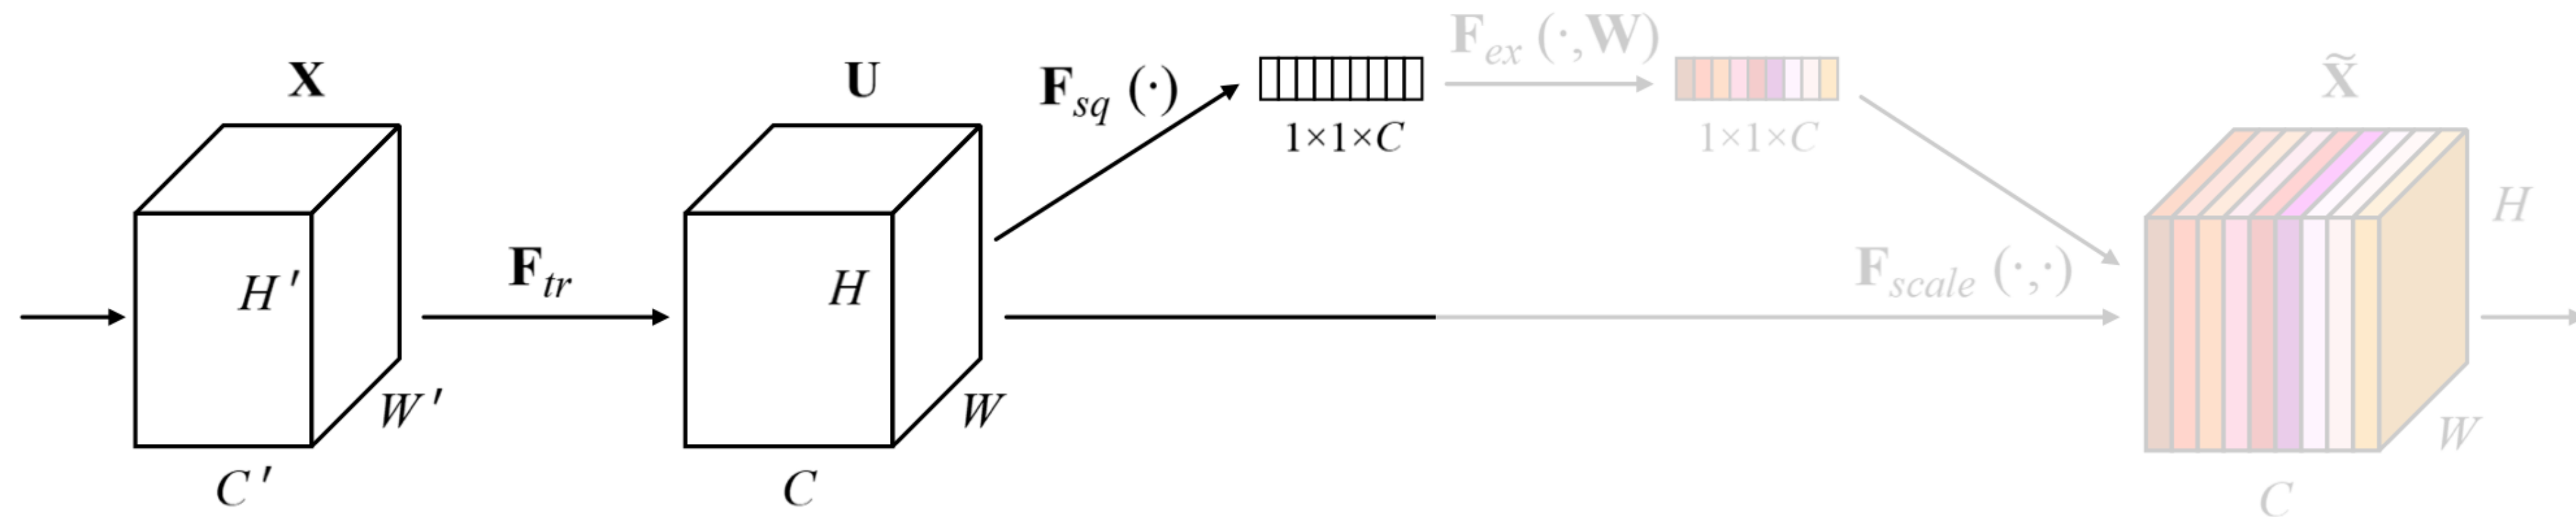
“Linear Bottlenecks”

Slightly less parameters than v1, same accuracy

Mark Sandler, Andrew Howard, Menglong Zhu, Andrey Zhmoginov, Liang-Chieh Chen,
MobileNetV2: Inverted Residuals and Linear Bottlenecks, <https://arxiv.org/abs/1801.04381>



SQUEEZE-AND-EXCITATION NETWORKS (SENET)



Intuition: model shall learn where to attend. Here: which channel

Also: number of channels often becomes larger with an increasing depth of the network

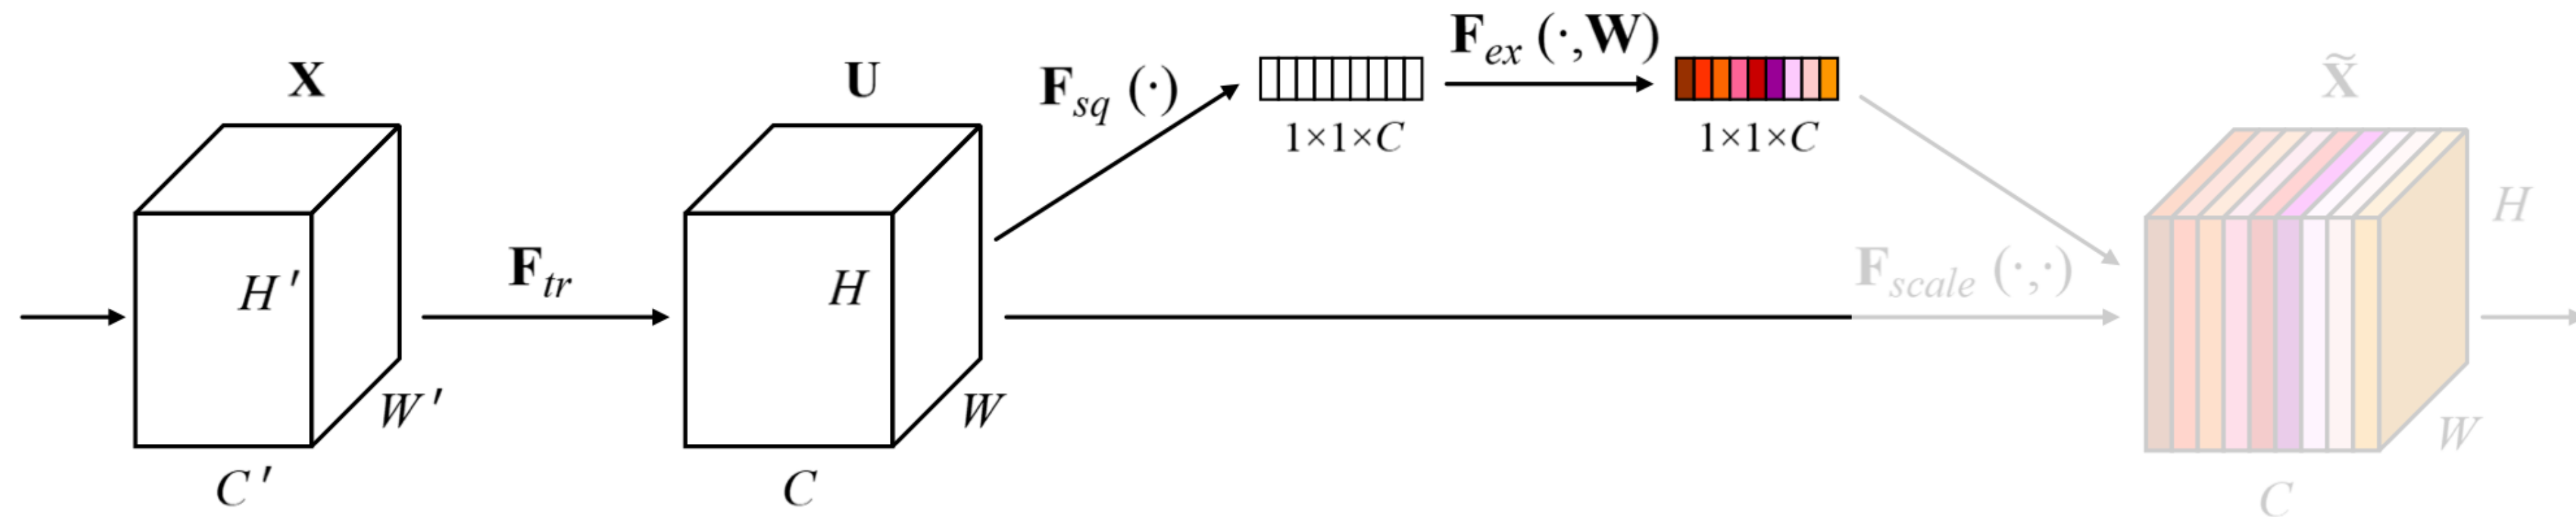
Squeeze module \mathbf{F}_{sq} to reduce information

Feature descriptor that decomposes in information of each feature map into a single value

Here by using global average pooling to transform an $H \times W \times C$ tensor into a $1 \times 1 \times C$ one

Essentially a vector of length C that encodes the feature descriptor of each fmap

SQUEEZE-AND-EXCITATION NETWORKS (SENET)



Excitation module \mathbf{F}_{ex} for adaptive recalibration

Fully capture channel-wise dependencies by learning

Not mutually exclusive: multiple channels can be important

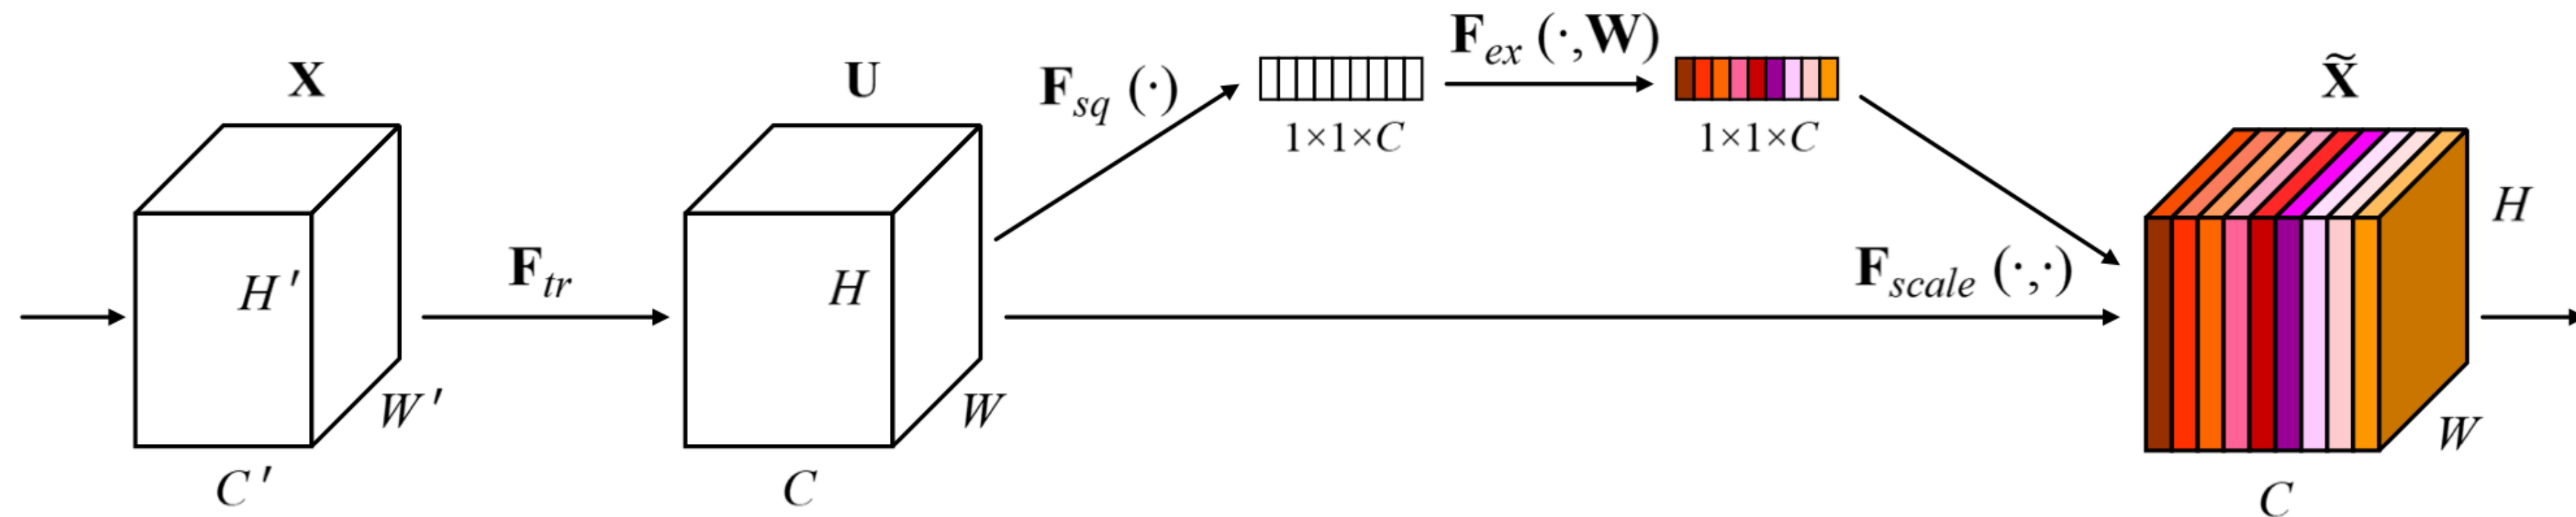
$$\Rightarrow \mathbf{F}_{ex}(\mathbf{z}, \mathbf{W}) = \sigma(\mathbf{W}_1 \cdot \text{ReLU}(\mathbf{W}_2 \cdot \mathbf{z}))$$

Dimensionality reduction layer (\mathbf{W}_2), followed by dimensionality-increasing layer (\mathbf{W}_1)

Shape is maintained, size of latent space can be controlled: C/r , for a reduction parameter r

Sigmoid activation σ scales the output values to a range $[0,1]$

SQUEEZE-AND-EXCITATION NETWORKS (SENET)



Scale module F_{scale} : channel-wise multiplication between the scalar output of F_{ex} and the corresponding fmap

Simple, computationally cheap

Downside?

Adds MACs and parameters

Scale is computationally cheap but adds plenty of state

MORE ARCHITECTURES

MNASNet (2018): use NAS to find efficient model architectures

Slightly outperforming MobileNet v2

But based on SE blocks that were explicitly made part of the search

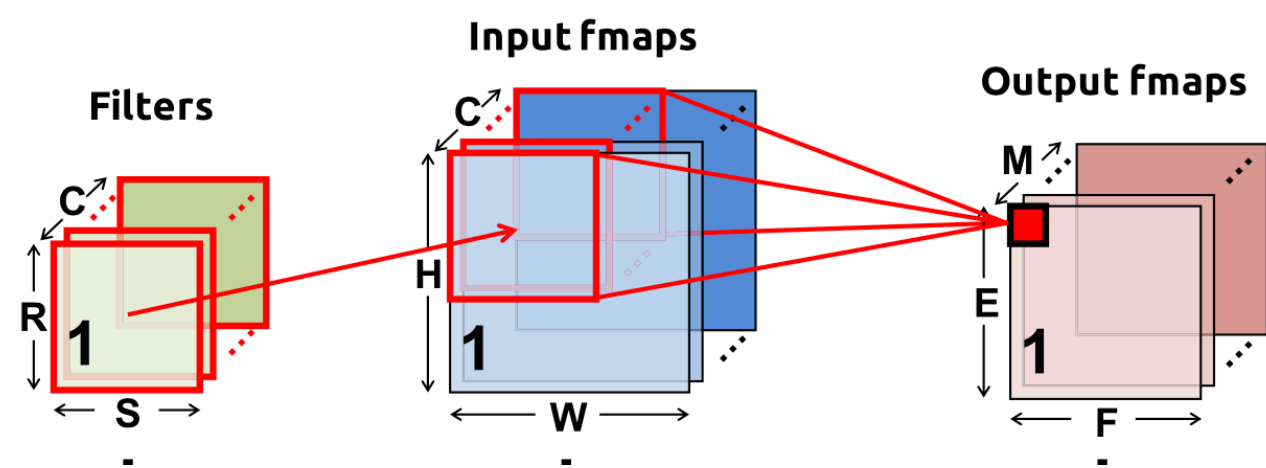
MobileNet v3 (2019): combining MNASNet and v2

EfficientNet (2019, 2020): combination of skip connections, depthwise-separable convolutions, SE blocks, and NAS

Optimization goal: $ACC(m) \cdot (FLOPs(m)/T)^w$

For test accuracy $ACC(m)$ and number of FLOPs $FLOPs(m)$ of a given model m , trade-off hyperparameter w that governs the relative importance, and target FLOP count T

Even more: SqueezeNext, LogicNets (FPGAs), ShuffleNet (2017, 2018), ...



EFFICIENCY METRICS

	MACs	Parameters (weight state)	Units (activation state)
FC	$MAC_f = WHCO$	$W_f = WHCO$	$U_f = O$
Convolution	$MAC_c = (EF \cdot RSC) \cdot M$	$W_c = RSCM$	$U_c = EFM$
Grouped convolution	$MAC_{cg} = \frac{MAC_c}{g}$	$W_{cg} = \frac{W_c}{g}$	$U_{cg} = EFM$
Depthwise separable convolution	$MAC_{cds} = EF \cdot (RSC + CM)$	$W_{cds} = RSC + CM$	$U_{cds} = EFC + EFM$

BEYOND CONVOLUTIONS

With material from Roger Grosse (U. Toronto, CSC421/2516)

RECURRENT NEURAL NETWORKS

Predicting sequences: speech-to-text and text-to-speech, caption generation, machine translation, signal processing

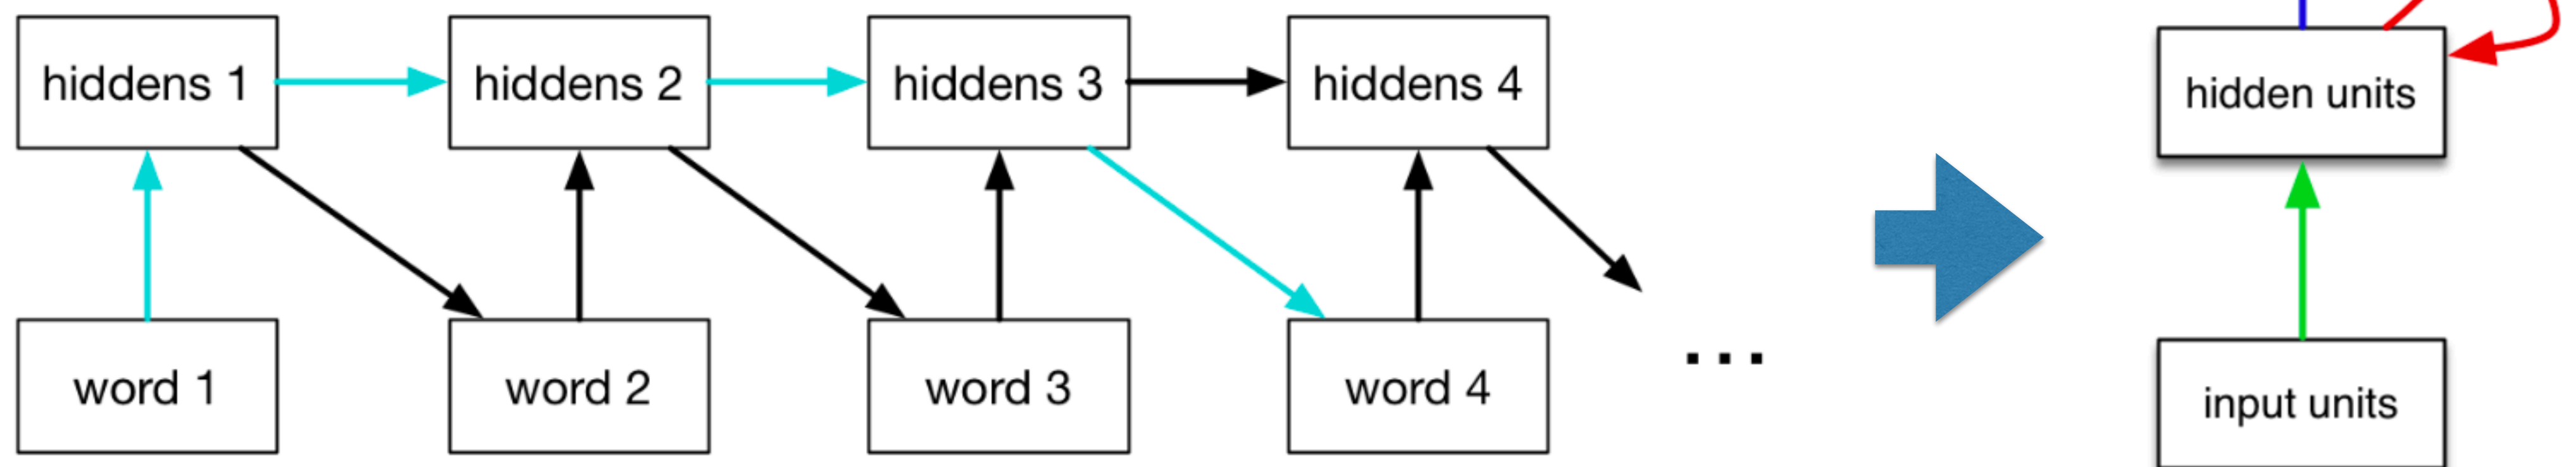
Sequence-to-sequence prediction: input is also a sequence

Recurrent neural networks (RNNs) have memory which is very helpful for sequences

For no memory, see neural language models and c.f. Markov assumption (model is memoryless)

=> RNNs are based on a hidden state/unit

Recurrent NN vs. feedforward NN



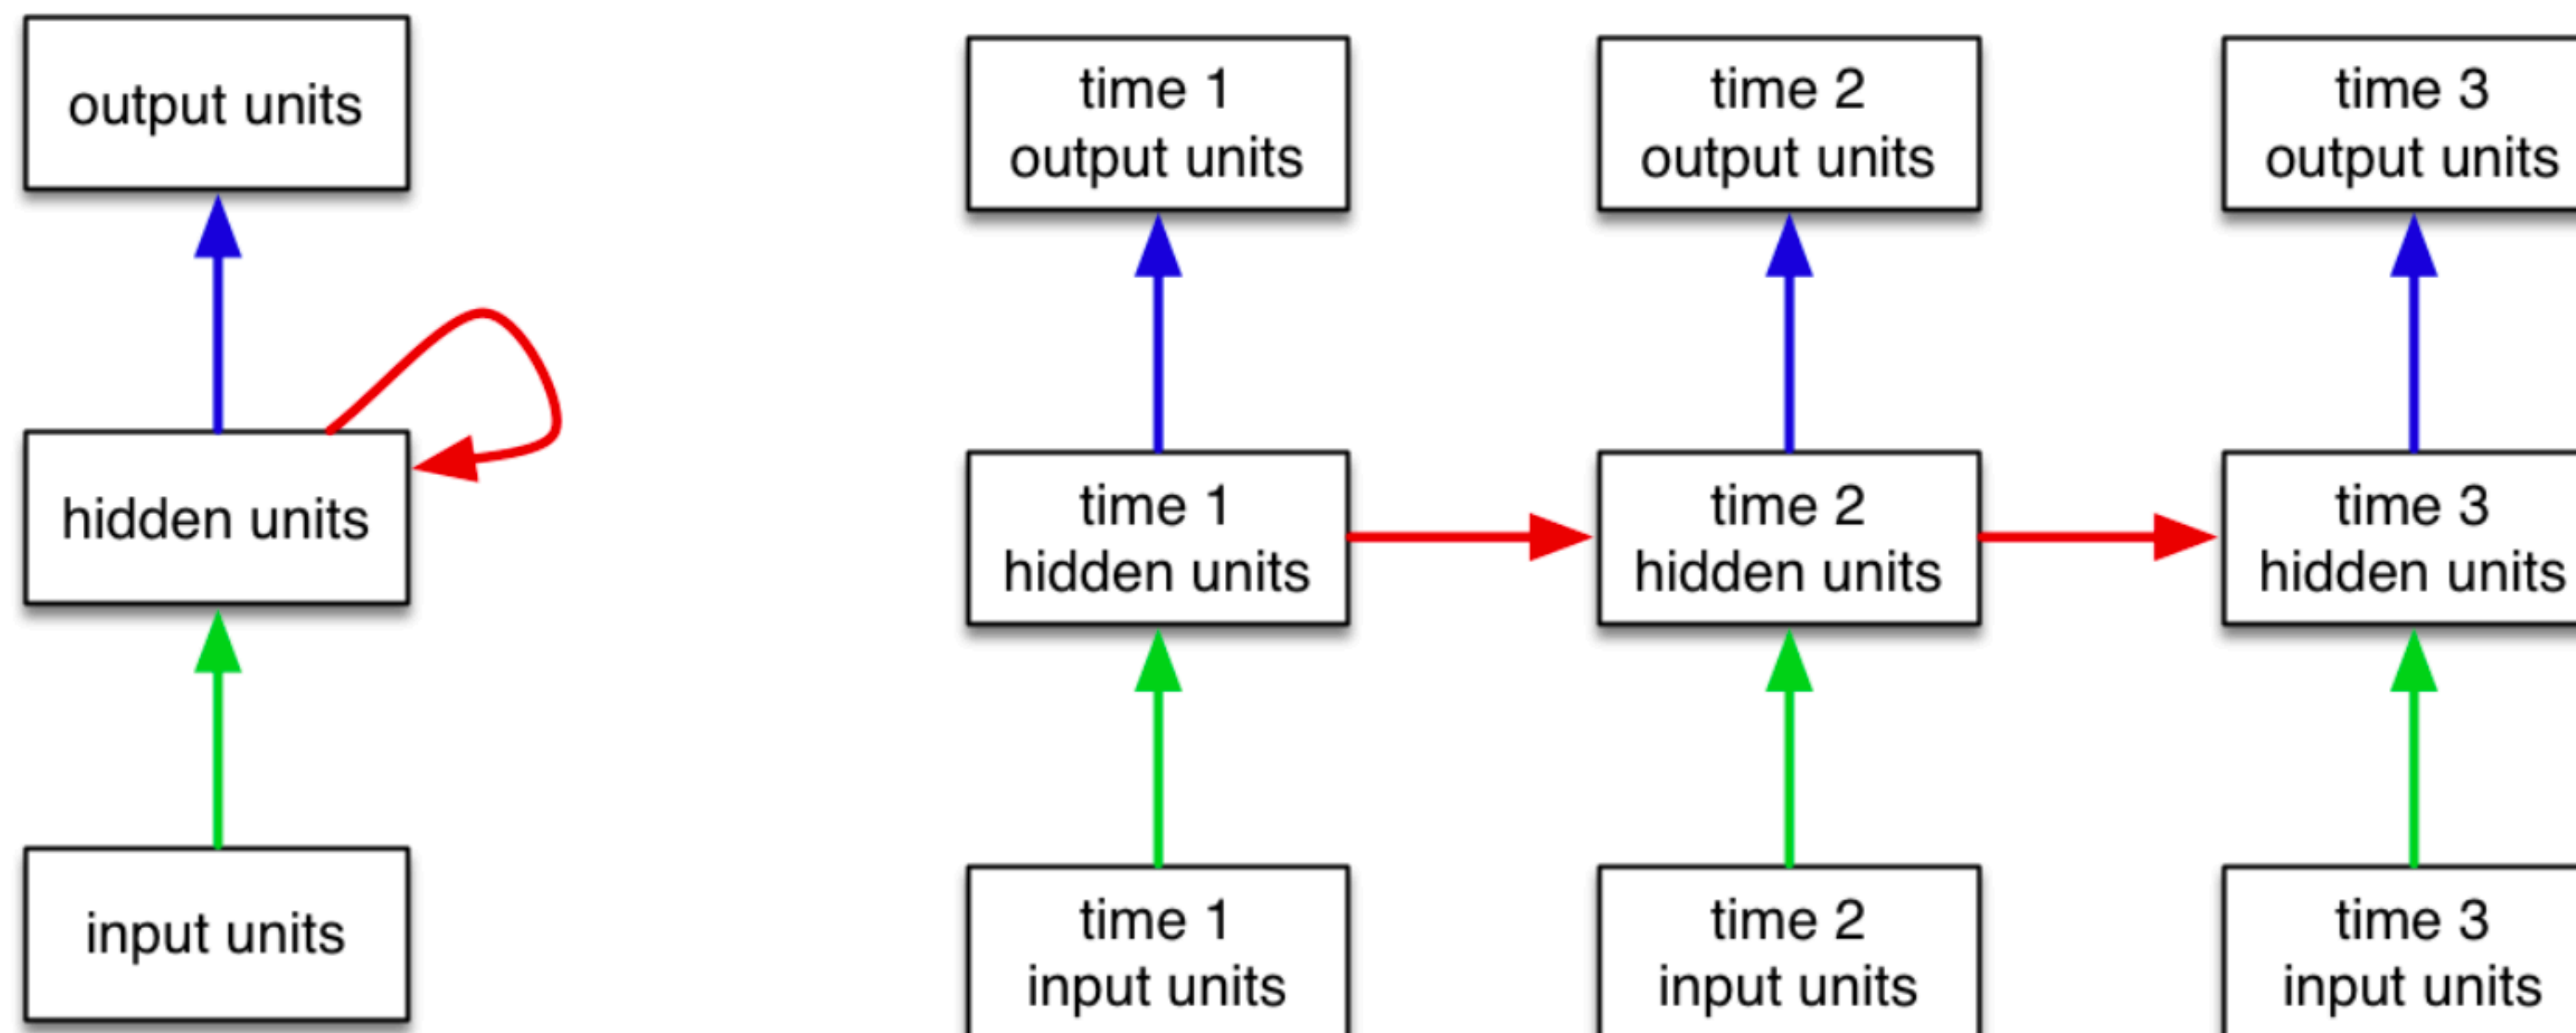
UNROLLING RNNS

RNN: dynamical system with one set of hidden units that feed into themselves

Unrolling: representing the units at all time steps

Parameters are shared among all time steps

Sometimes the biases for the first time step are treated differently



BACKPROP THROUGH TIME

Unrolling allows to re-use backpropagation to train RNNs

Careful with weight sharing (u , w , v)

Unrolling effectively increases the depth of the model => backprop is very sensitive to vanishing and exploding gradients

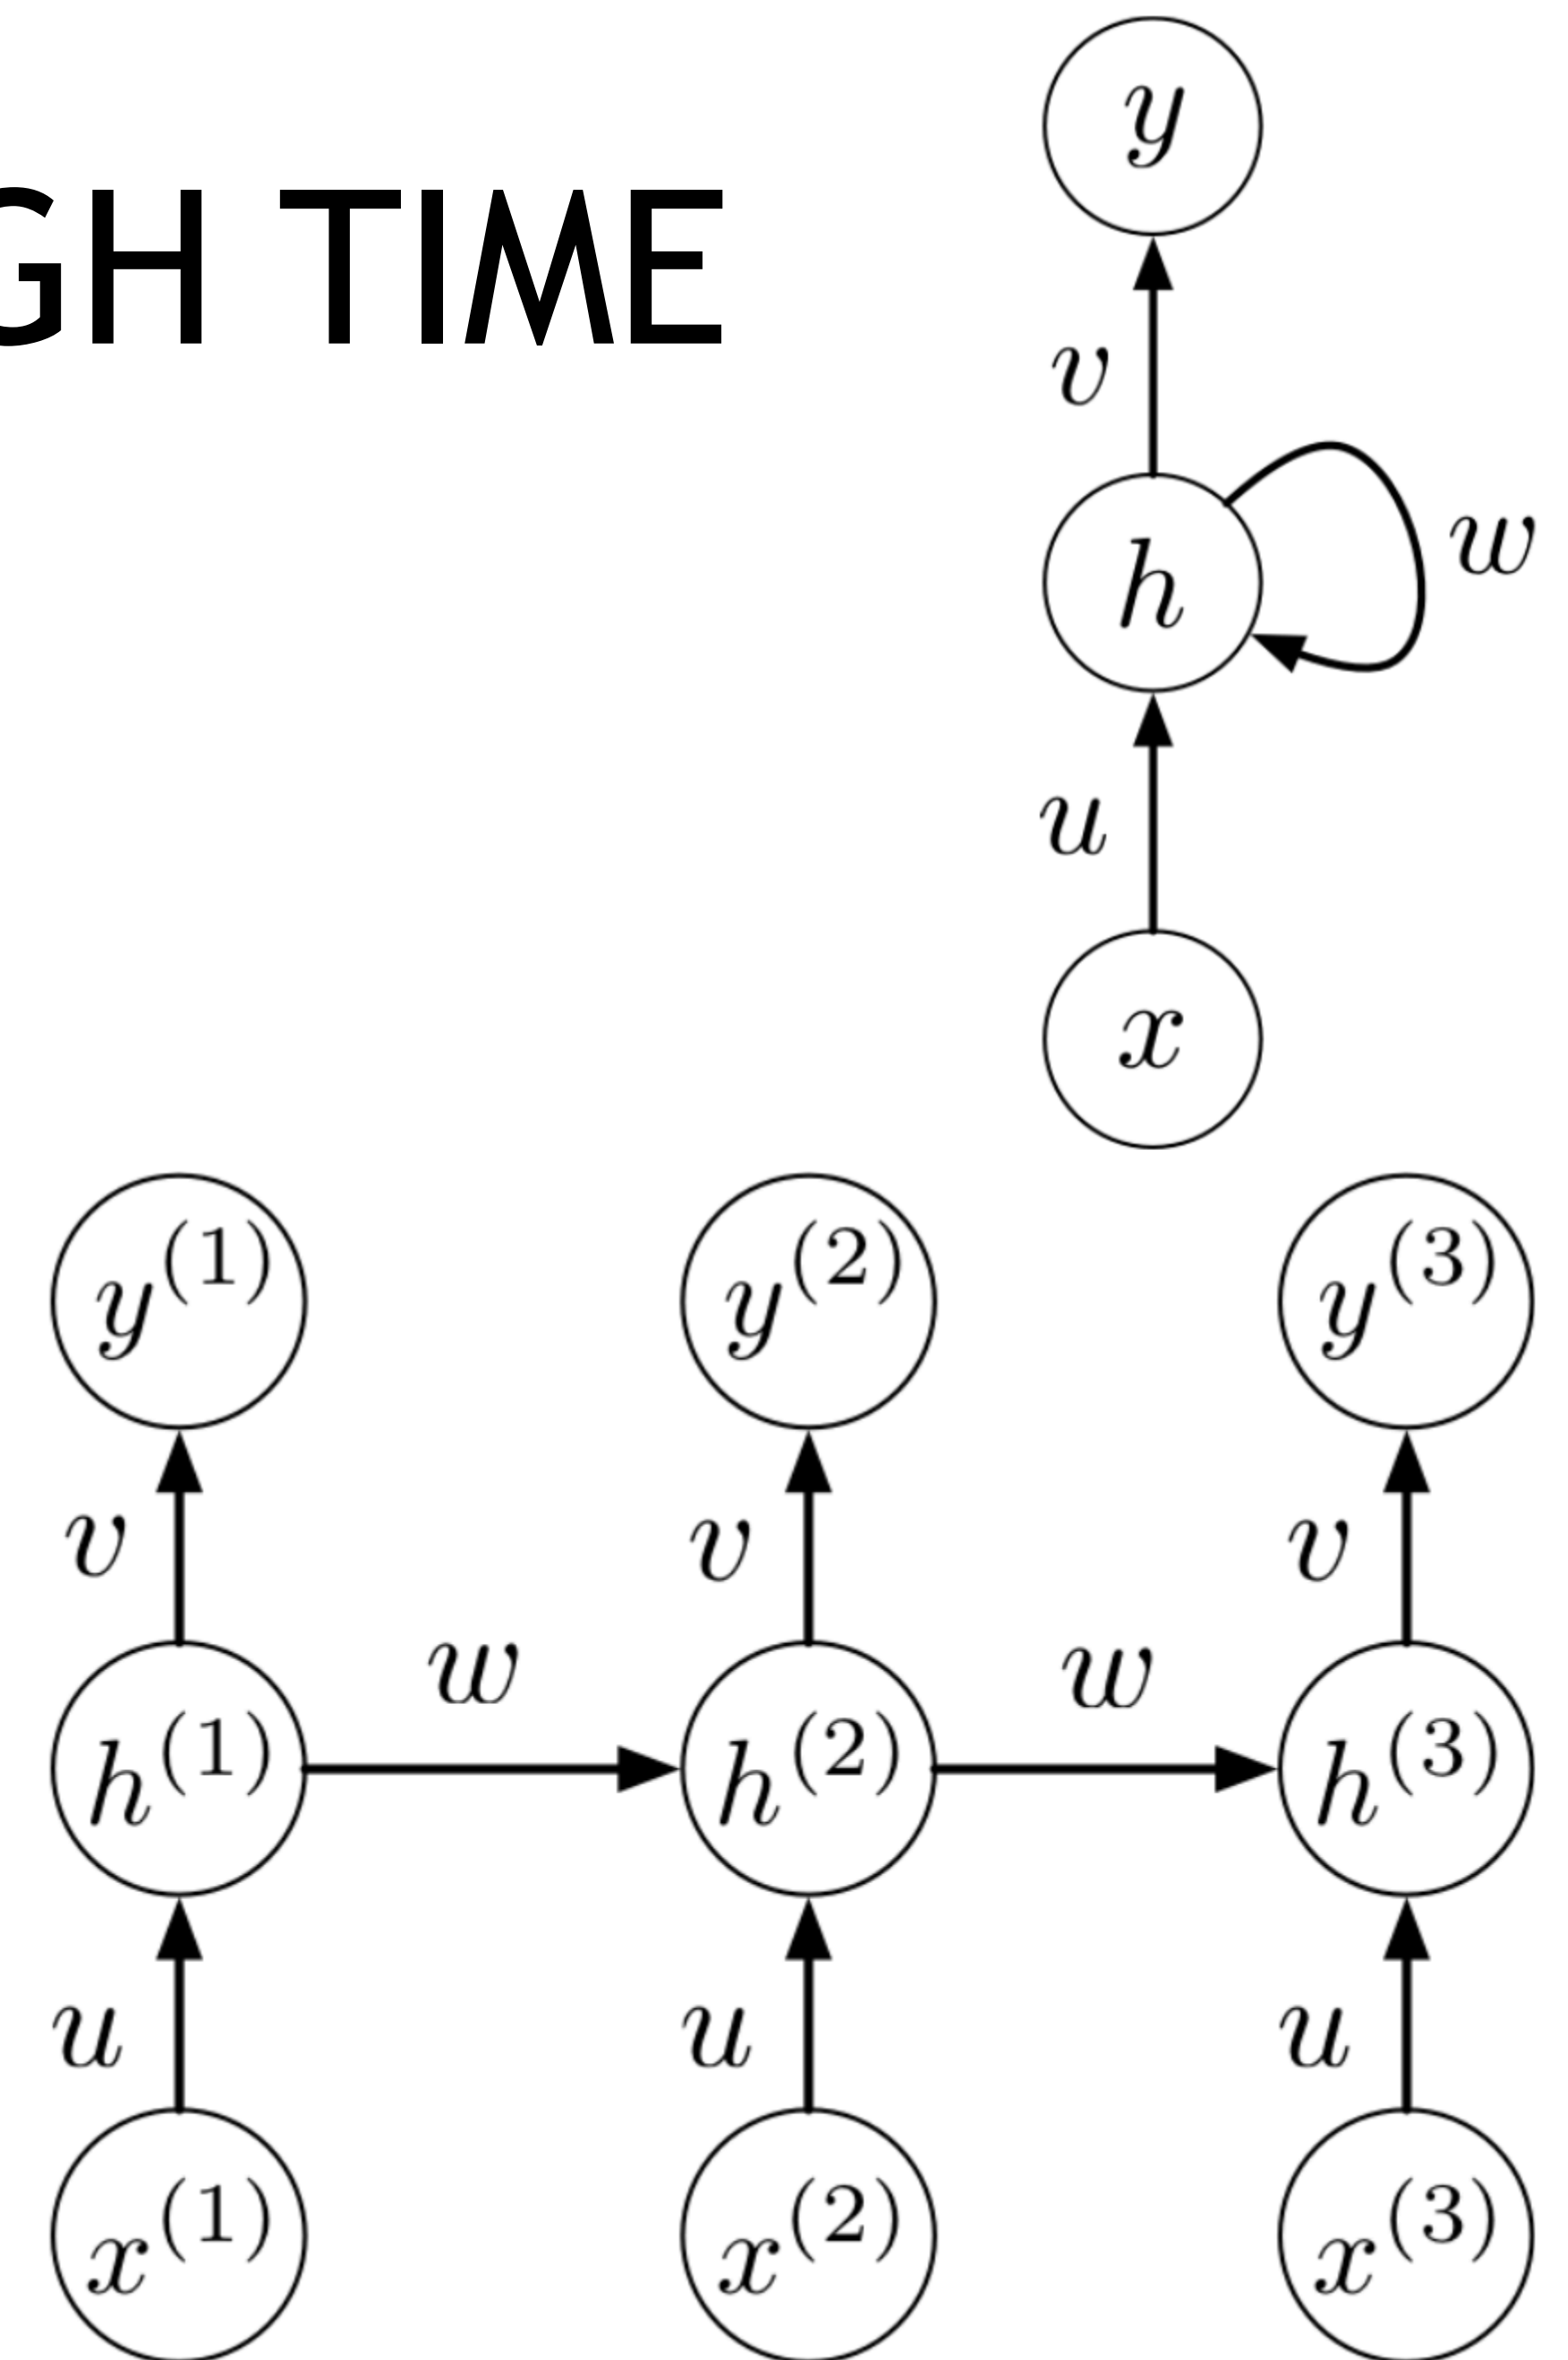
=> Essentially simple, very difficult in practice

Usually RNNs are small

Not a prime example for model compression

Contrary, inherently resource-efficient

Other scalability issues, as difficult to predict long sequences



ATTENTION FOR LONG-RANGE SEQUENCE MODELING

Attention-based Neural Machine Translation (NMT) encoder-decoder

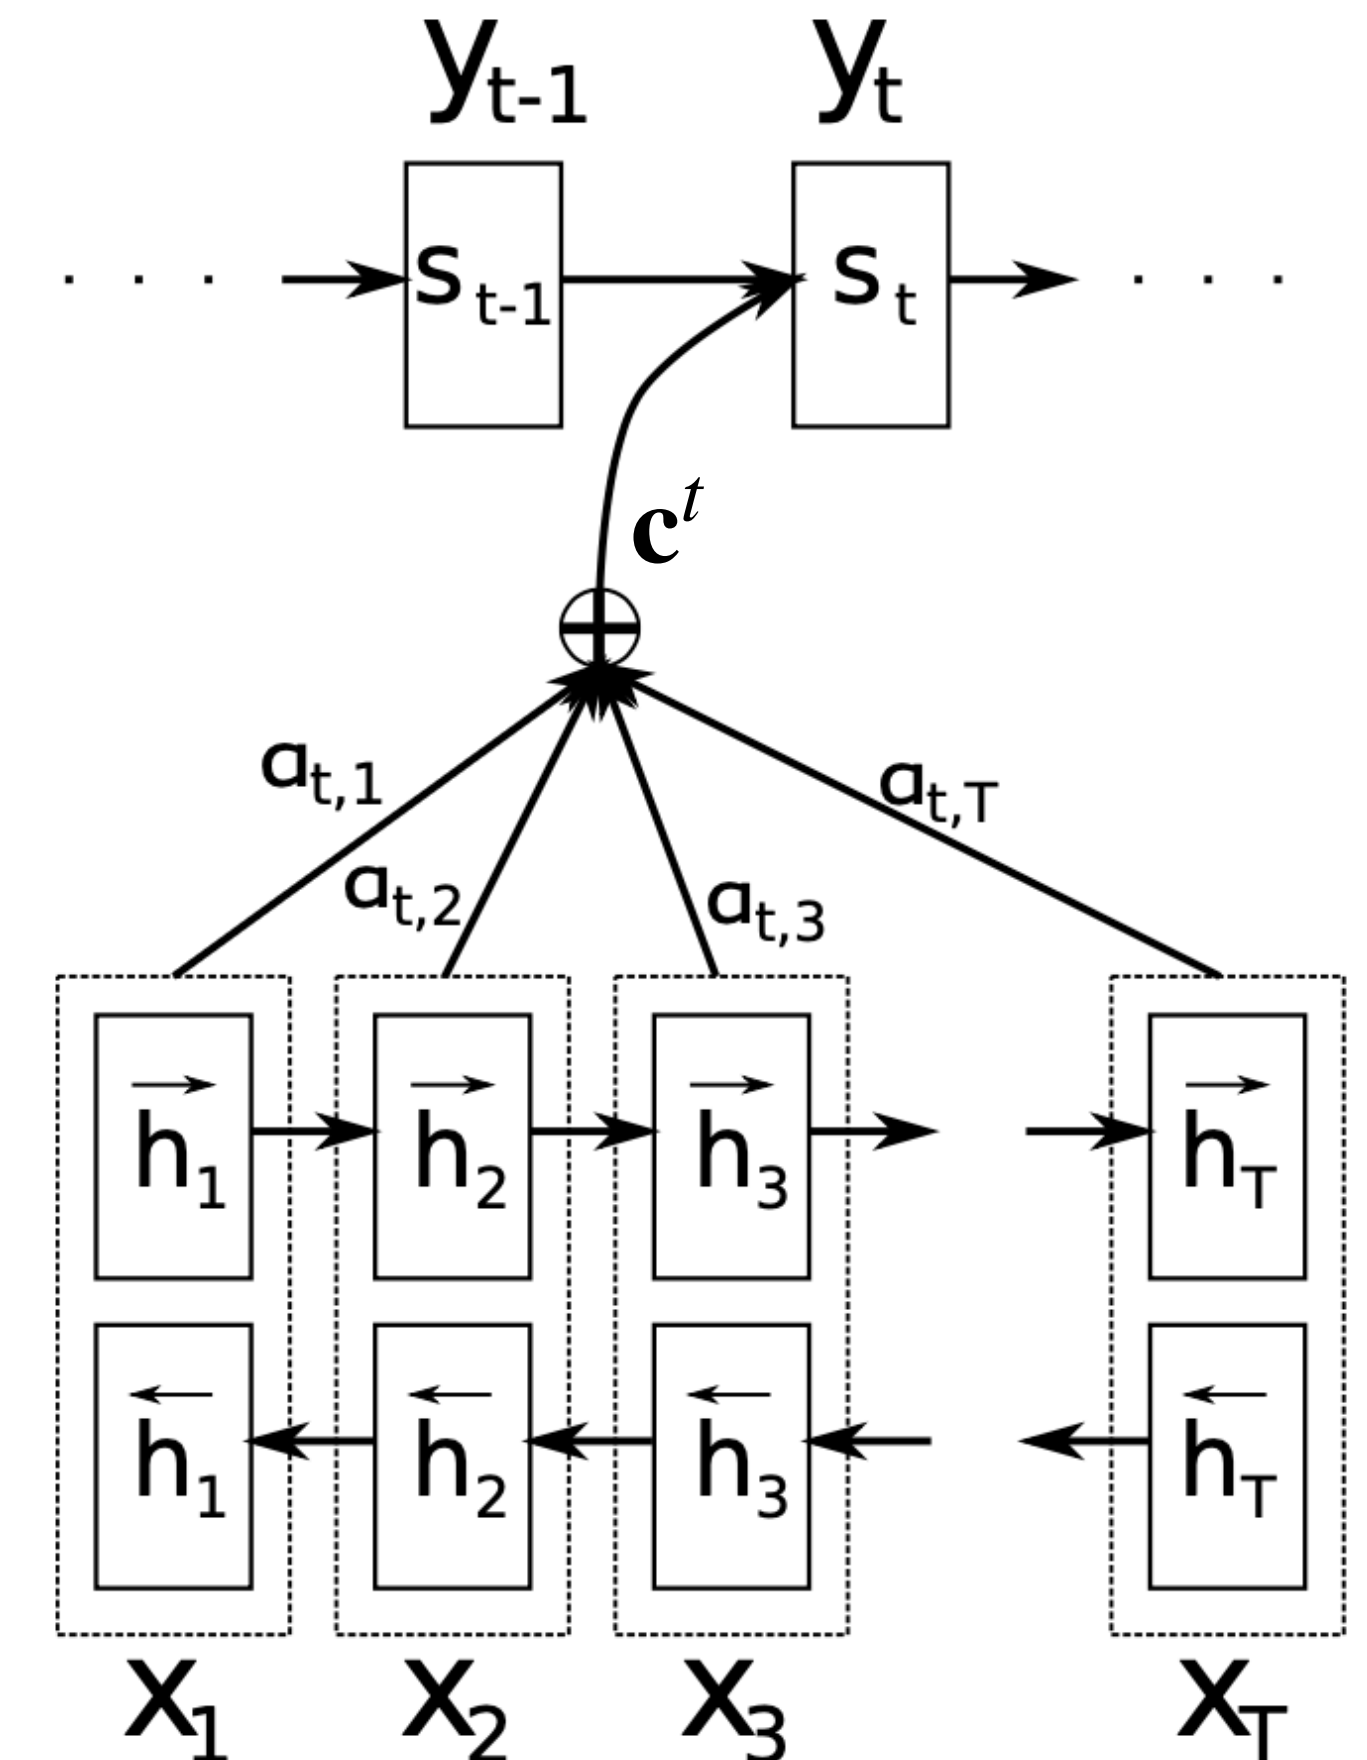
1. Encode the source sentence \mathbf{x} (e.g., French) into a sequence of hidden states
2. Attend selectively to different parts of the source at each decoding step.
3. Decode by generating the target sentence \mathbf{y} (e.g., English) one token at a time

Bidirectional RNN as encoder, normal RNN as decoder

Encoder: Forward hidden states $\vec{\mathbf{h}}$, and backward hidden states $\overleftarrow{\mathbf{h}}$
(annotation vector \mathbf{h} by concatenating $\vec{\mathbf{h}}$ and $\overleftarrow{\mathbf{h}}$)

Attention: compute context vector \mathbf{c}^t for each time step t

Decoder: decides which part of the source to pay attention to, and combines with own hidden state \mathbf{s}



ATTENTION

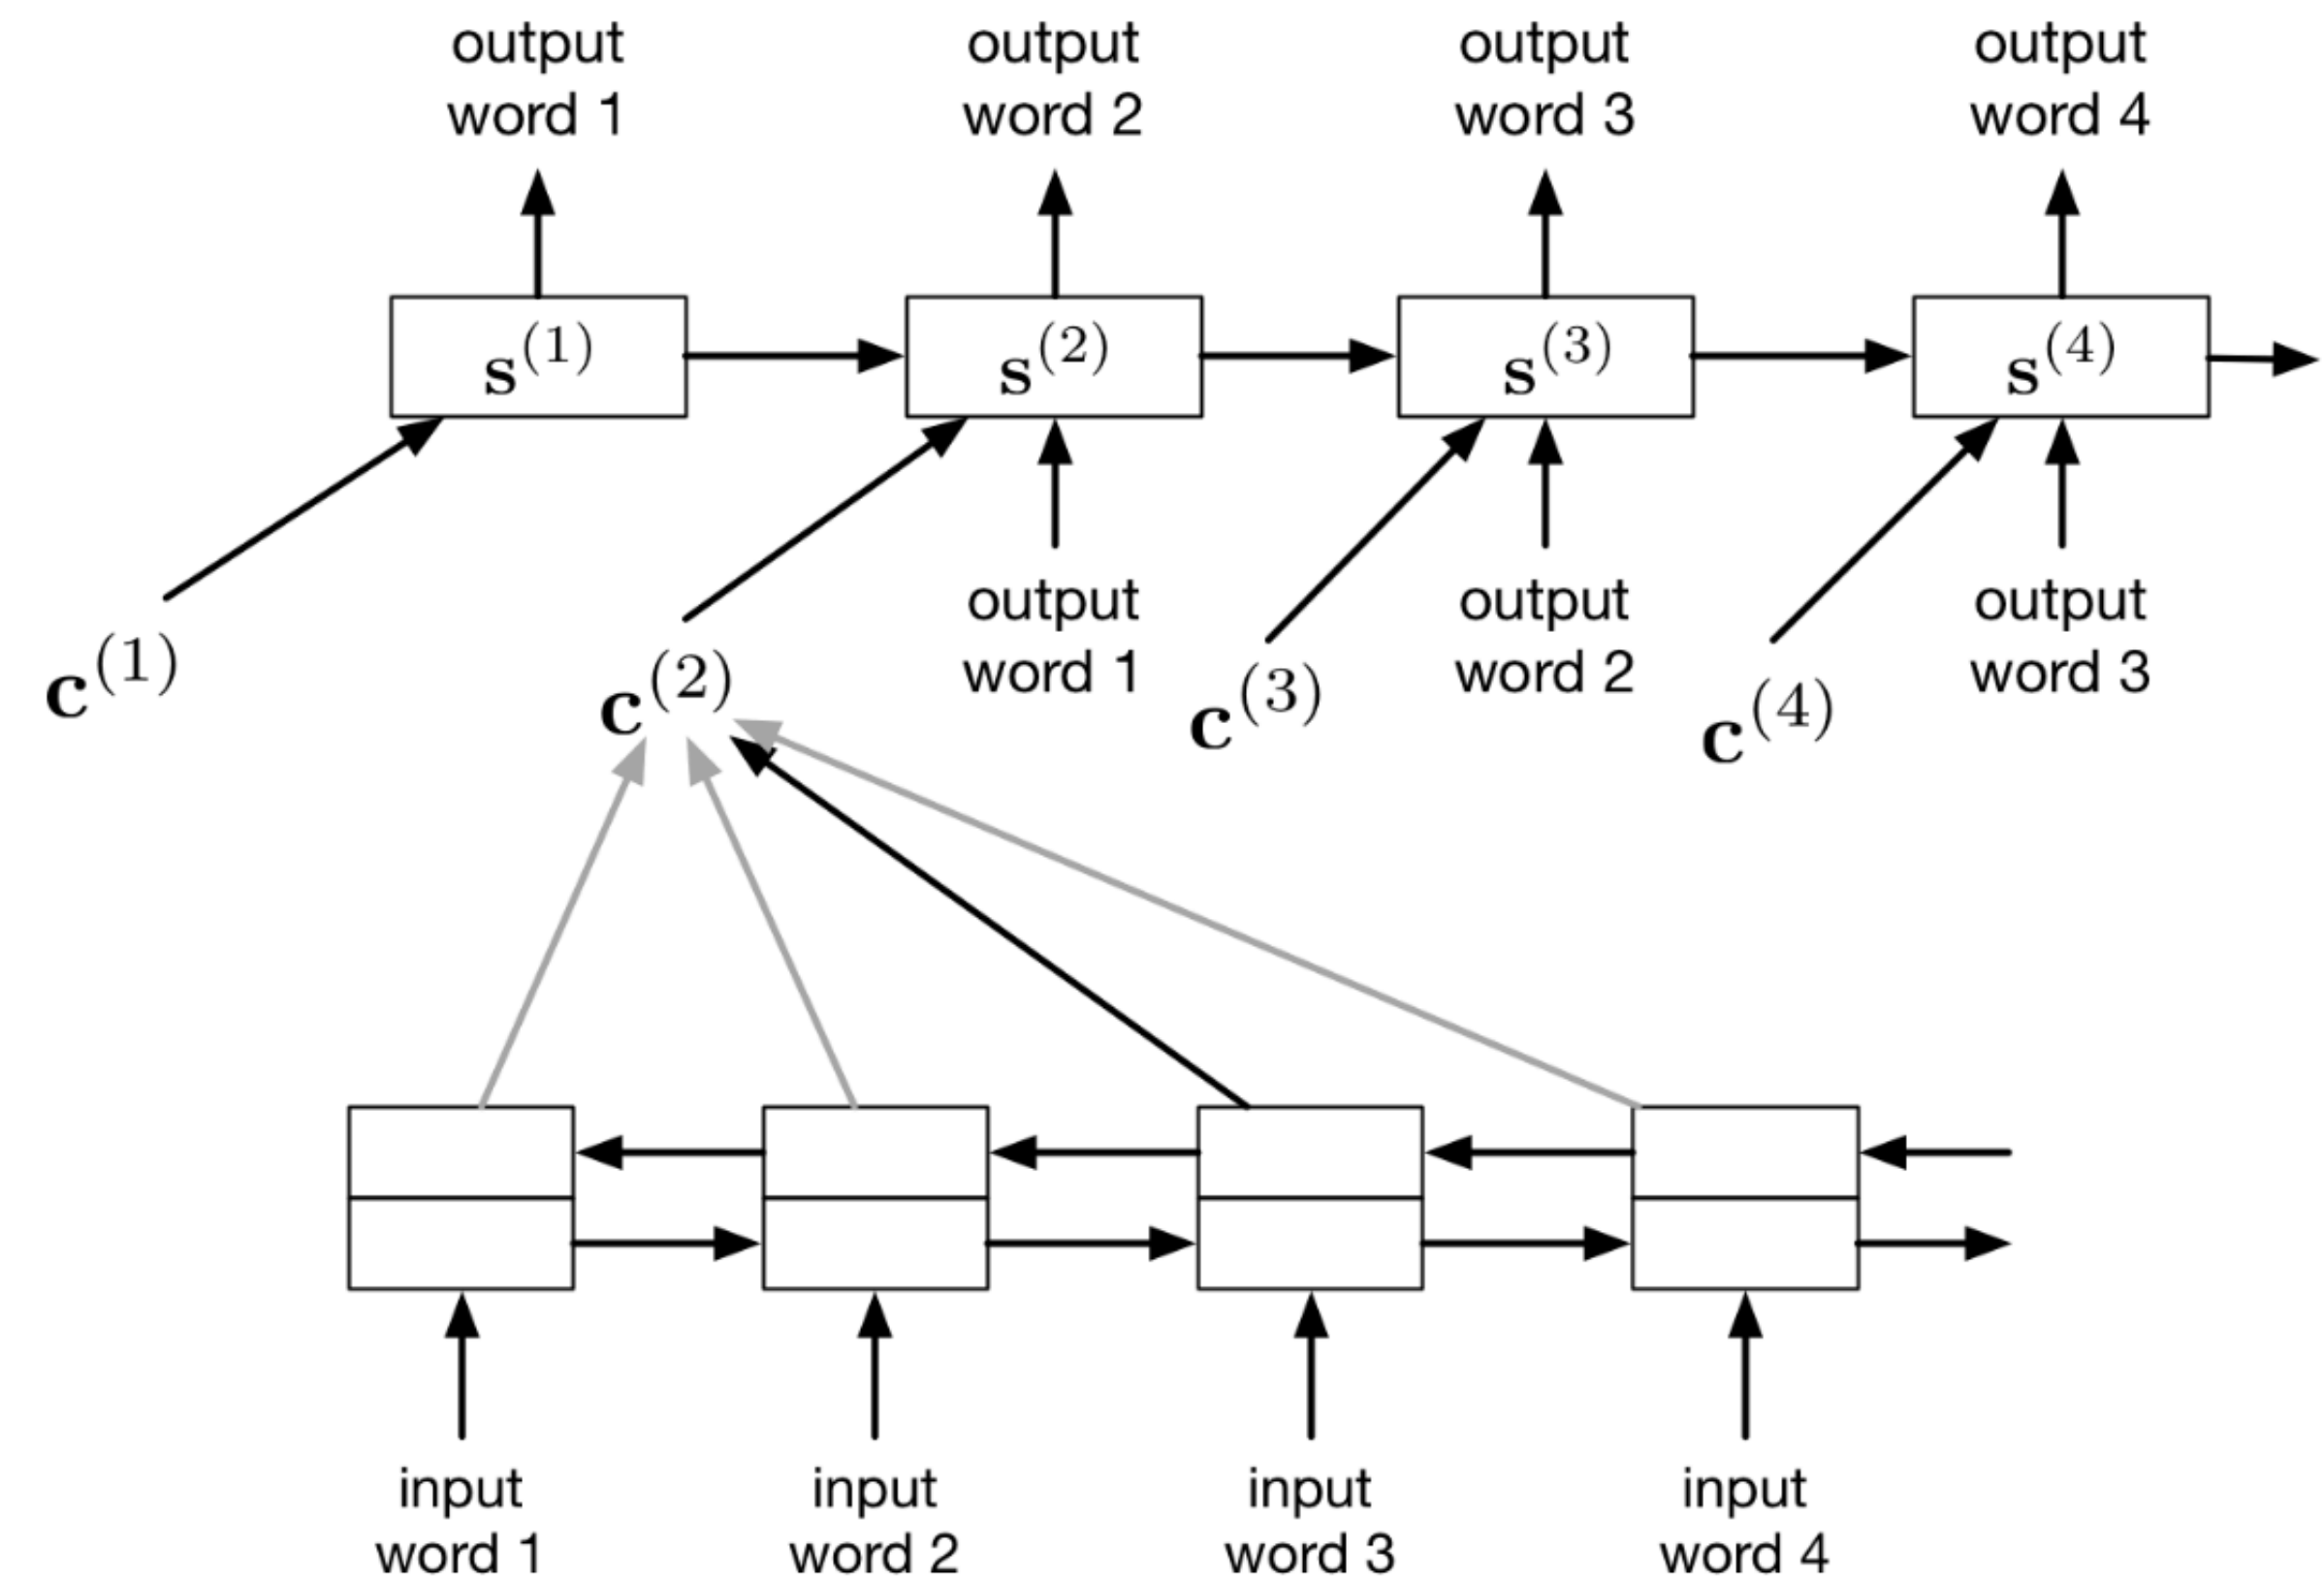
Context is a weighted sum of the encoder annotations: $c^i = \sum_j \alpha_{ij} \cdot h_j$

Context vector needs to be computed every time step

Weight α_{ij} of each annotation h_j

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_k \exp(e_{ik})}$$

, for an alignment model a that scores how well input at position j and output at position i match: $e_{ij} = a(s_{i-1}, h_j)$



ATTENTION

Content-based addressing: attention function depends on annotation vector rather than on position in the sentence

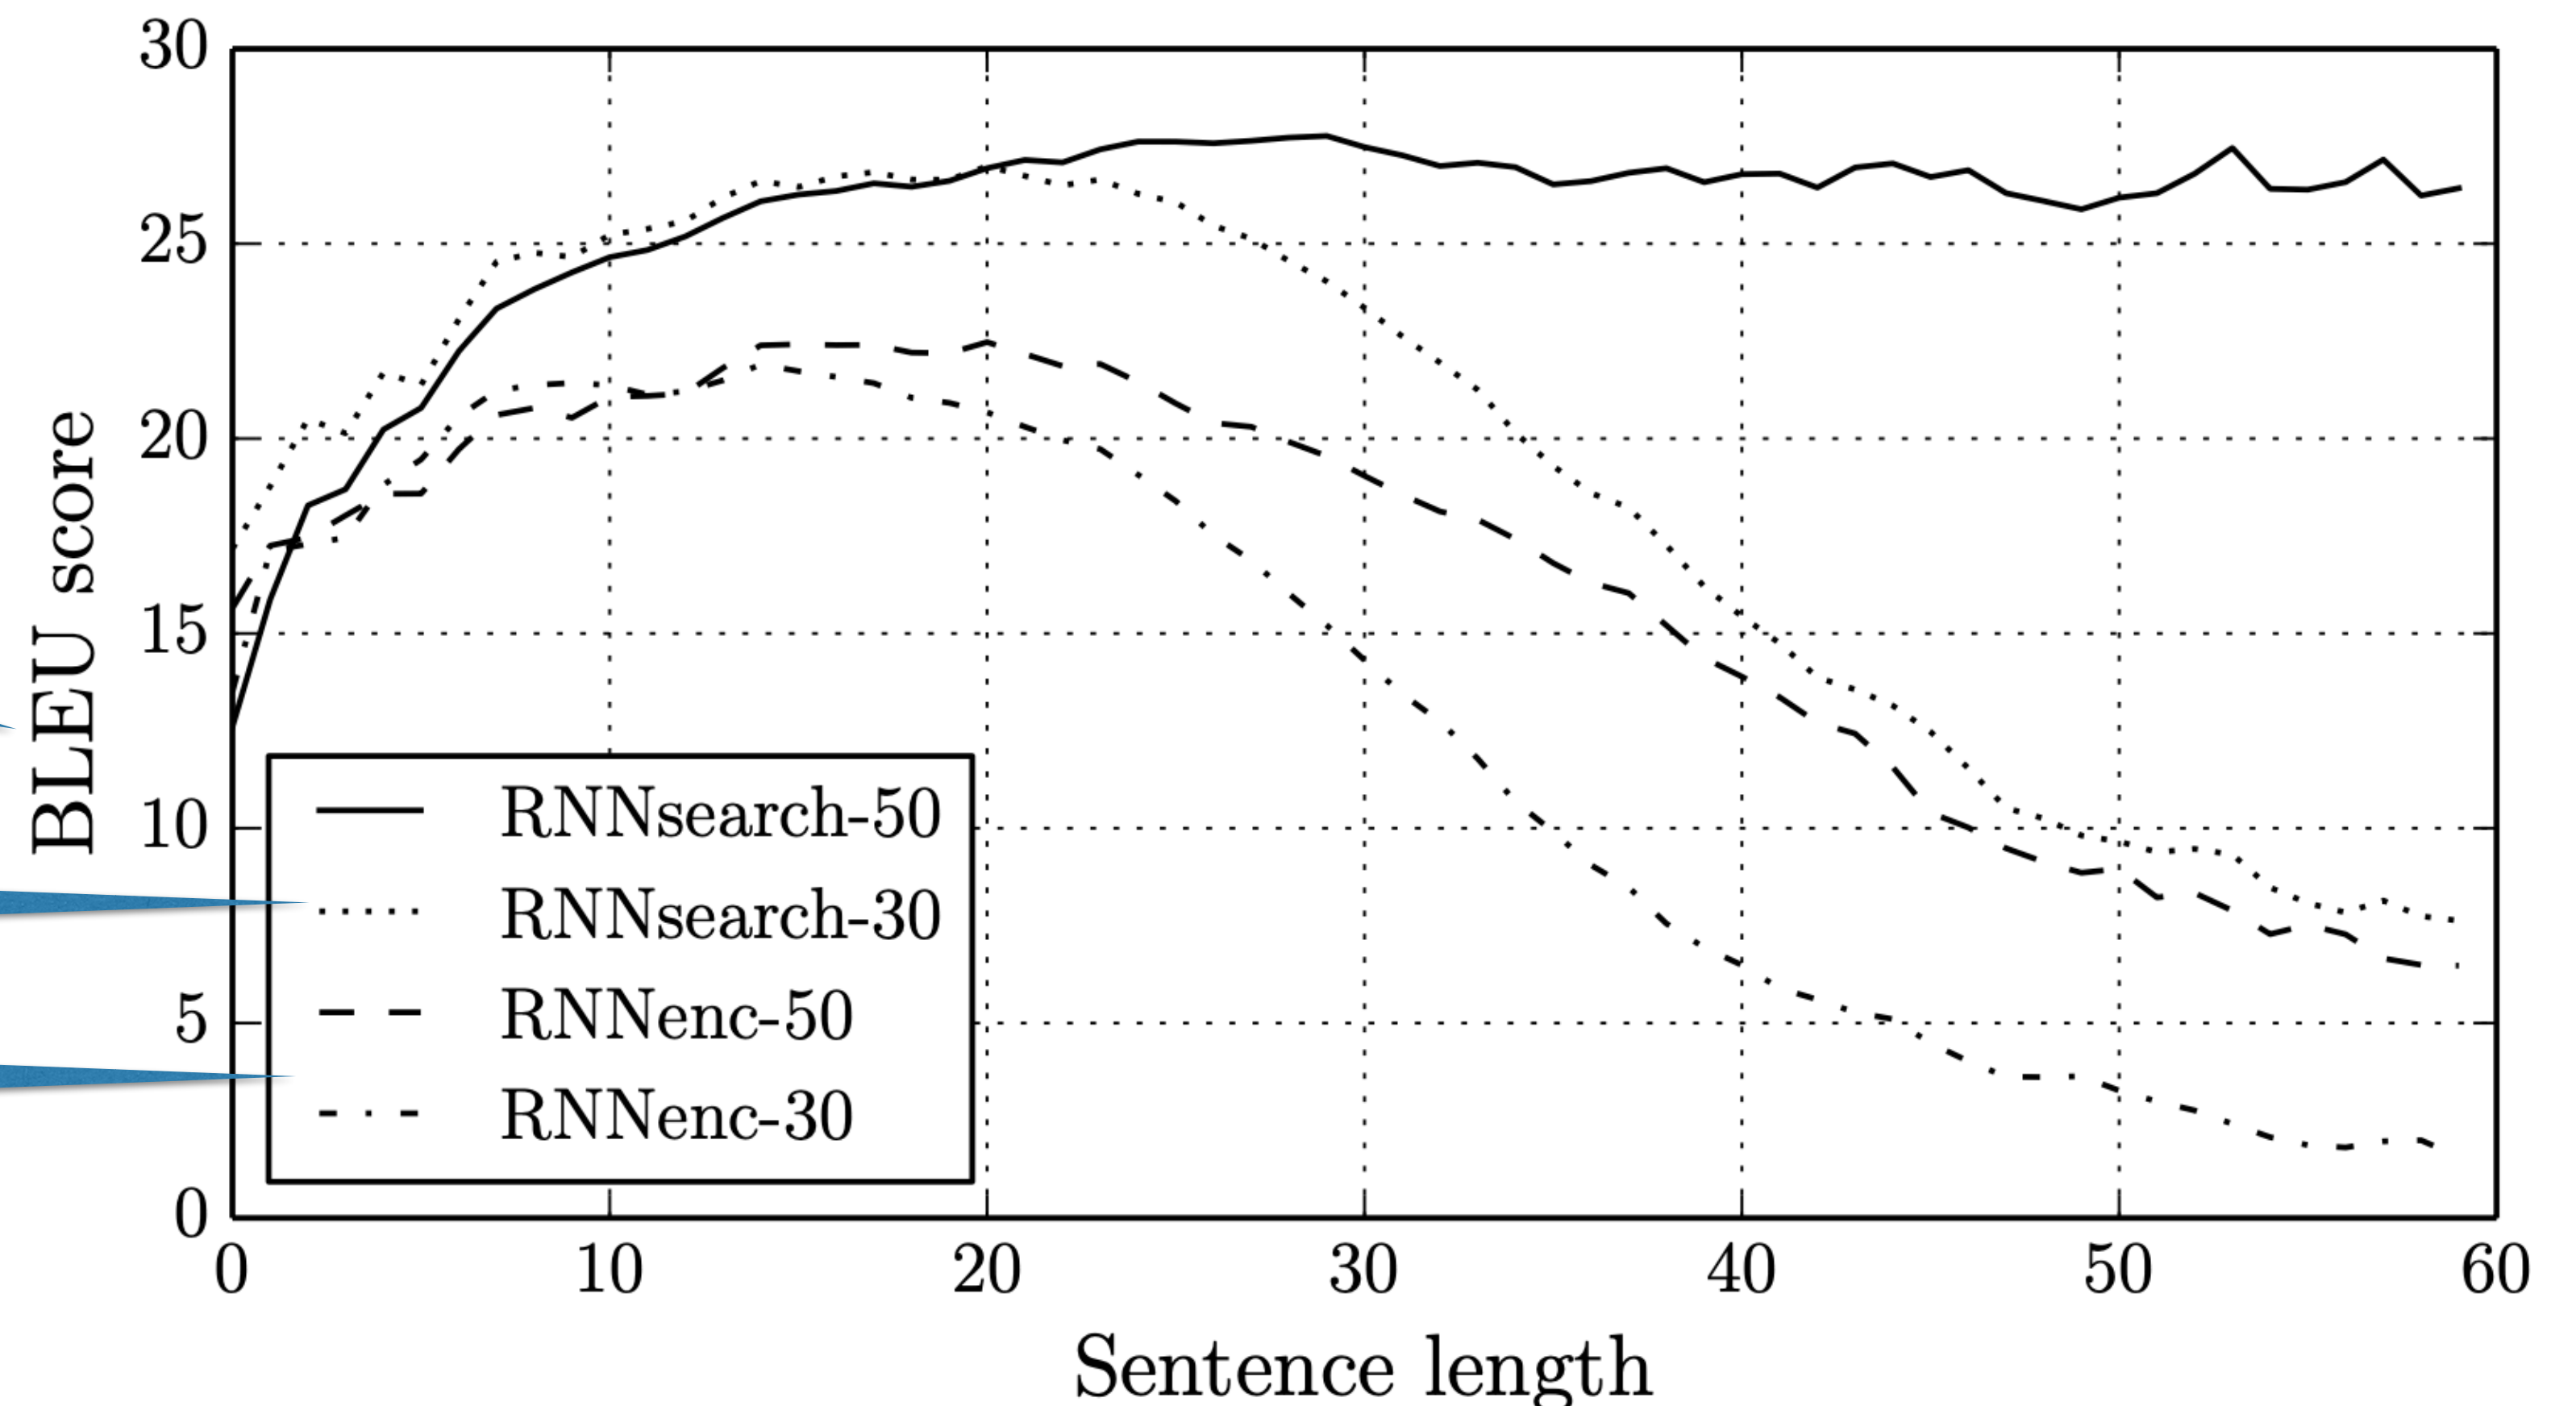
Inline with “biology”

C.f. convolutions

BLEU score: higher is better

Attention-based encoder-decoder

Plain RNN encoder-decoder



COMPUTATIONAL COSTS

Sequence length t , number of layers d , number of neurons per layer k

Transformer is not covered here, but essentially it replaces RNNs with feedforward model architectures

	Training complexity	Training memory	Test complexity	Test memory
RNN	tk^2d	tkd	tk^2d	kd
RNN + attention	t^2k^2d	t^2kd	t^2k^2d	tkd
Transformer	t^2kd	tkd	t^2kd	tkd

WRAPPING UP

SUMMARY

Many advanced tricks to improve resource efficiency of convolutional neural networks

Pay attention on “efficiency metrics” (FLOPs != latency, state is not only defined by number of parameters, ...)

There is more than MACs and parameters

RNNs extend feed-forward models by having memory

But are difficult to train, effectively reducing the scalability of such architectures

Attention mechanisms is much more effective than plain memory (all history)

Transformers proven to be much more scalable, now also used in image processing, document processing, etc.

Image classification foundational for many image processing tasks

Natural language processing continues this story, but at a different scale

=> Era of foundational models

